



Introduction to Machine Learning



Wintersession 2024
Jan 16-18, 22-23

Gage DeZoort
Julian Gold
Jake Snell



With materials from:

Brian Arnold, Jonathan Halverson, Christina Peters, Savannah Thias, and Amy Winecoff

Course Outline

Date	Instructor	Topic
Day 1 Tue. 1/16	Jake Snell	Machine Learning Overview and Simple Models
Day 2 Wed. 1/17	Jake Snell	Model Evaluation and Improving Performance
Day 3 Thu. 1/18	Julian Gold	Introduction to Neural Networks
Day 4 Mon. 1/22	Gage DeZoort	Survey of Neural Network Architectures
Day 5 Tue. 1/23	Gage DeZoort Julian Gold Jake Snell	Hackathon! <ul style="list-style-type: none">• Computer Vision• Diffusion Models• Large Language Models

Agenda

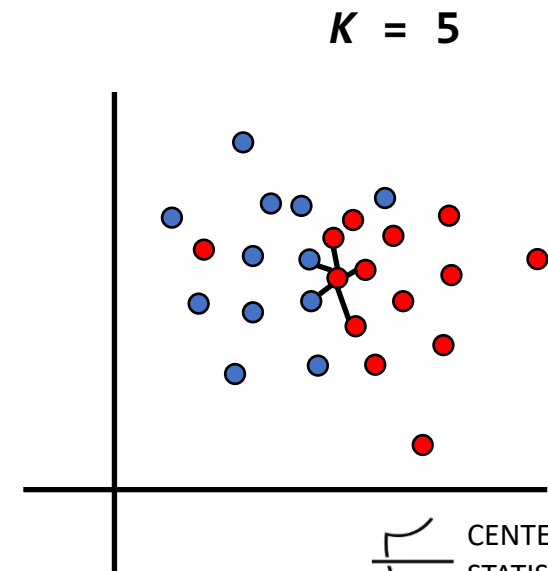


- K-nearest Neighbors
 - Regression and classification
- Clustering with K-means
- Evaluation paradigms and improving performance

Intro to K -nearest neighbors (KNN)

- simple but powerful
- can be used for classification *or* regression!
- algorithm
 1. for a given test sample (yellow dot), find the K nearest training samples in feature space
 - 2a. for **classification**, assign label by majority vote
 - 2b. for **regression**, assign value by mean of neighbors

K is a tunable parameter!

- choose value that gives better predictions on test data



Coding in Python!

https://github.com/PrincetonUniversity/intro_machine_learning/tree/main/day2
<https://jdh4.github.io/intro-ml>

Google Colab

1. Open notebook (.ipynb):




2. Click “Open in Colab”:



JupyterLite

1. Download notebook (.ipynb) from left.
2. Open JupyterLite:



3. Upload the notebook and open:



K-Means Clustering Visualization

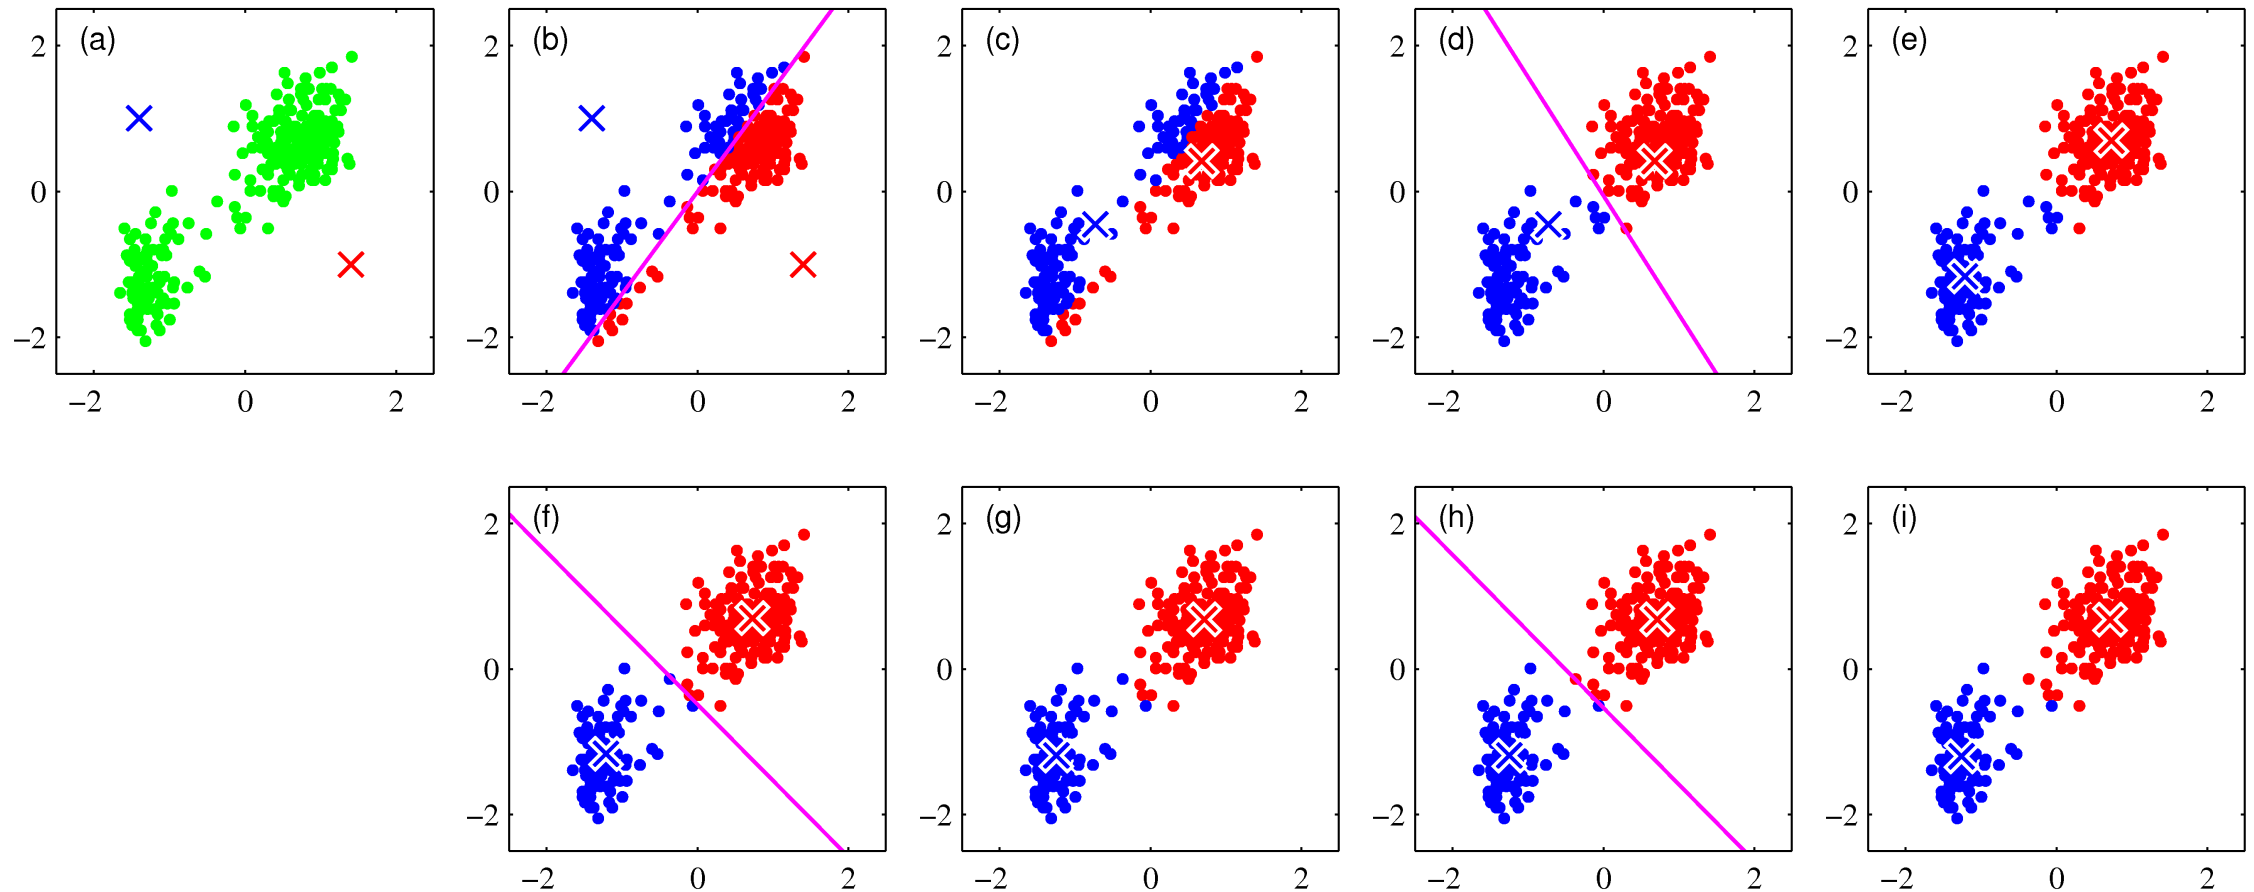


Figure credit: Bishop, Christopher M. 2006. *Pattern Recognition and Machine Learning*.

Overview of Machine Learning Process

1. Define the problem to be solved.
2. Split the data into train / validation / test.
3. Run the validation loop:
 - a. Choose a set of models.
 - b. Train each model by optimizing its parameters on the training set.
 - c. Evaluate the performance of each model on the validation set.
 - d. Repeat until performance is satisfactory.
4. Evaluate final performance on the test set.

Datasets? Input features?
Targets? Evaluation metrics?

