



A Hands-On Introduction to Machine Learning



Wintersession 2025
January 15–17, 21

Julian Gold
Gage DeZoort



With materials from:

Brian Arnold, Gage DeZoort, Julian Gold, Jonathan Halverson, Christina Peters, Savannah Thias, Amy Winecoff

Mini-Course Outline

Date	Topic	Instructor
Wed. 1/15	Machine Learning Overview and Simple Models	Julian Gold
Thu. 1/16	Model Evaluation and Improving Performance	Julian Gold
Fri. 1/17	Introduction to Neural Networks	Gage DeZoort
Tue. 1/21	Survey of Neural Network Architectures	Gage DeZoort
Wed.+Thu. 1/22-1/23	Getting Started with LLMs with PLI	Simon Park, Abhishek Panigrahi
Wed. 1/22	Graph Neural Networks for Your Research	Gage DeZoort
Wed. 1/22	Machine Learning for the Physical Sciences	C. Jespersen, R. Pastrana, Q. Gallagher, H. Johnson

Artificial Intelligence


A science devoted to making machines think and act like humans.

Machine Learning

Focuses on enabling computers to perform tasks without explicit programming.

Deep Learning

A subset of machine learning based on artificial neural networks.



WHY USE DL?

- Data with complex (highly non-linear) relationships
- Big data – many examples to leverage
- High-dimensional data
- Data with complicated structure (images, video, language, social networks etc.)

20 DEEP LEARNING Applications




Deep neural network

Input layer


Multiple hidden layers

Output layer



DEEP NEURAL NETWORKS

- A class of ML algorithms based on *artificial neural networks* (ANNs)
- Neural network → networks of neurons responding to stimuli
- “Deep” → multiple layers of neurons interacting in sequence



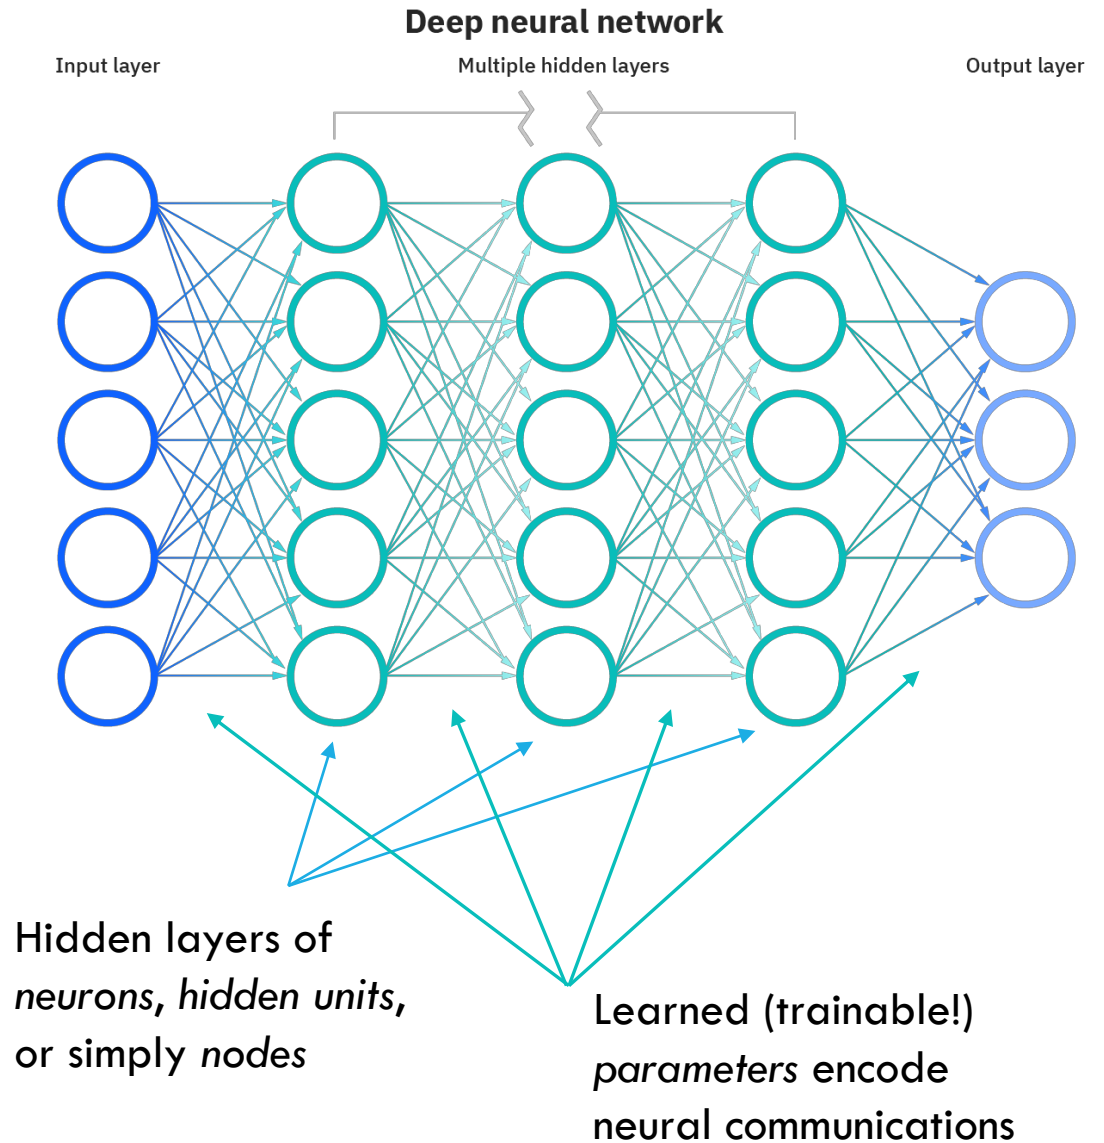
Hidden layers of
neurons, hidden units,
or simply *nodes*

DEEP NEURAL NETWORKS

- A class of ML algorithms based on *artificial neural networks* (ANNs)
 - Neural network → networks of neurons responding to stimuli
 - “Deep” → multiple layers of neurons interacting in sequence
- Practically speaking, DNNs are non-linear models designed to leverage complicated relationships in data

$$f(x \mid \text{parameters}) = \text{output}$$


Adjustable, e.g. fit to data
in supervised learning



SIMPLE NN


- **Example:** Classification problem with $x_i \in \mathbb{R}^2$ $y_i \in \{0,1\}$
- Let's draw a test data point and pass it through a simple ANN:

$$x = (x_1, x_2) \quad y = 1$$



- Single hidden layer with 2 neurons
- **Key Ingredients:**
 - Trainable weights $W_1, W_2, W_3, W_4, W_5, W_6$ and biases b_1, b_2, b_3
 - Non-linear activation functions (called non-linearities) $\sigma(z)$

SIMPLE NN



1. Compute *preactivations* at each neuron

$$z_1 = w_1x_1 + w_2x_2 + b_1$$

$$z_2 = w_3x_1 + w_4x_2 + b_2$$

Or, in matrix notation:


$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

the *weight matrix* mixes up the inputs and feeds them to the each neuron

the *bias vector* adds constants to the mixed inputs

Note: this is a linear operation!!

SIMPLE NN



1. Compute *preactivations* at each neuron

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

2. Calculate how much the neuron *activates* given the strength of the *pre-activation*

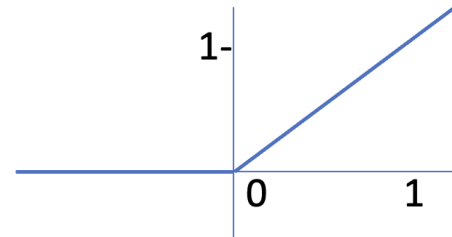
$\sigma(z) \rightarrow$ activation function

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$


NON-LINEAR ACTIVATION FUNCTIONS

- Non-linear activation functions allow us to learn complex relationships by “intervening” between linear operations
- In practice, they allow neurons to switch “on” and “off” to varying degrees


Rectified Linear Unit (ReLU)



Hyperbolic Tangent

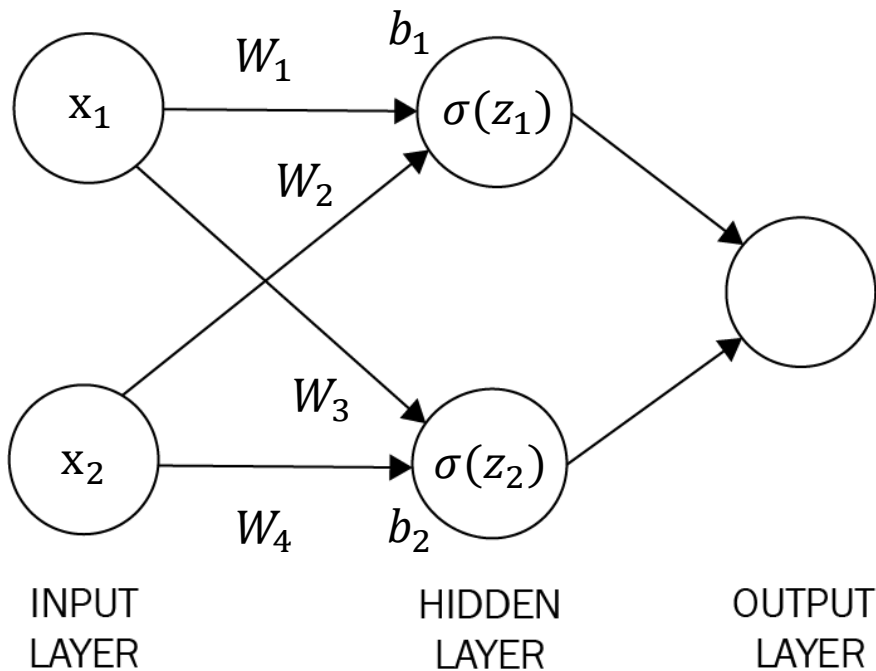


Sigmoid



← Most popular choice; simple to compute (fast training), no “saturating” regions with tiny gradients

SIMPLE NN




1. Compute *preactivations* at each neuron

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

2. Calculate how much the neuron *activates* given the strength of the *pre-activation*

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$

HOW IT WORKS



1. Compute *preactivations* at each neuron

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$


2. Calculate how much the neuron *activates* given the strength of the *pre-activation*

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$

3. Calculate model outputs

$$\hat{y} = \begin{pmatrix} w_5 & w_6 \end{pmatrix} \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix} + b_3$$

HOW IT WORKS




Forward Pass → Predictions

1. Compute *preactivations* at each neuron
2. Calculate how much the neuron *activates* given the strength of the *pre-activation* (Repeat 1 and 2 to some fixed depth)
3. Calculate model outputs

When we define an NN, we fix an *architecture* by specifying:

- Number of hidden layers (here 1)
- Dimension of the hidden layers (here 2)
- Activation functions
- Random initial values for weights and biases

TRAINING A NN



The loss function may be very complicated in practice!

Training a NN \rightarrow find “optimal” weights, biases

- Start by defining a **loss function** $L(y, \hat{y})$
- Compute the gradient of $L(y, \hat{y})$ with respect to each weight and bias
- Update the weights and biases via **gradient descent**:


$$W_1^{(k+1)} = W_1^{(k)} - \gamma \frac{\partial L}{\partial W_1} \Big|_{W_1^{(k)}}$$

$$b_1^{(k+1)} = b_1^{(k)} - \gamma \frac{\partial L}{\partial b_1} \Big|_{b_1^{(k)}}$$

etc.

$\gamma \rightarrow$ learning rate

LEARNING RATES




TRAINING A NN

Backward Pass (Backpropagation)

Apply the chain rule to calculate derivatives of the loss with respect to the weights/biases

What we want: $\frac{\partial L}{\partial w}, \frac{\partial L}{\partial b}$

Generically, “gradients”



The diagram illustrates a vertical stack of four circular nodes representing the layers of a neural network. From bottom to top, the nodes are labeled w , x , y , and z . Upward-pointing arrows connect the nodes, indicating the forward pass. To the left of the arrows, the activation functions are labeled: $f^{(1)}$ between w and x , $f^{(2)}$ between x and y , and $f^{(3)}$ between y and z .

$$\begin{aligned}\frac{\partial z}{\partial w} &= \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w} \\ &= f^{(3)'}(y) f^{(2)'}(x) f^{(1)'}(w)\end{aligned}$$

PSEUDOCODE


```
for each training epoch:  
    for each (input, truth) in train_data:  
        prediction = NN(train_data)  
        loss = loss_function(prediction, truth)  
        gradients = compute_gradients(loss)  
        new_params = grad_descent(gradients, NN.parameters, lr)  
        model.update(new_params)
```

Deep neural network

Input layer

Multiple hidden layers

Output layer



SHOULD YOU USE DL FOR YOUR PROJECT?

Why DL?


- Produce predictions in multiple feed-forward stages → powerful feature extraction:
 - Less need for data pre-processing
 - Ability to leverage large amounts of data (“big data”)
- Handle more complicated data representations like images, sentences, and graphs
- Can be designed to perform complicated (multi-stage) tasks end-to-end

Why not DL?


- Developing and training large DL models is computationally expensive (slow, resource intensive) and often requires specialized computing hardware
- It usually takes a lot of data to train DL models

MORE ARCHITECTURES

- NNs are the building blocks for more complicated architectures, e.g.
- **Convolutional Neural Networks** are frequently applied to images or other grid data
- **Recurrent Neural Networks** are applied to sequences like sentences
- **Graph Neural Networks** operate on networks of objects (nodes) connected by their relationships (edges)
- **Generative Adversarial Networks** are used to generate new data (e.g. photographs) similar to a reference set




[Convolutional Neural Network: A Step By Step Guide | by Shashikant | Towards Data Science](#)



[Towards understanding glasses with graph neural networks \(deepmind.com\)](#)

Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org



A bit outdated, but fun to see the creativity...

[The mostly complete chart of Neural Networks, explained | by Andrew Tch | Towards Data Science](#)



Time for some NN practice!

Please navigate to Day 3!

[PrincetonUniversity/intro machine learning \(github.com\)](https://github.com/PrincetonUniversity/intro_machine_learning)