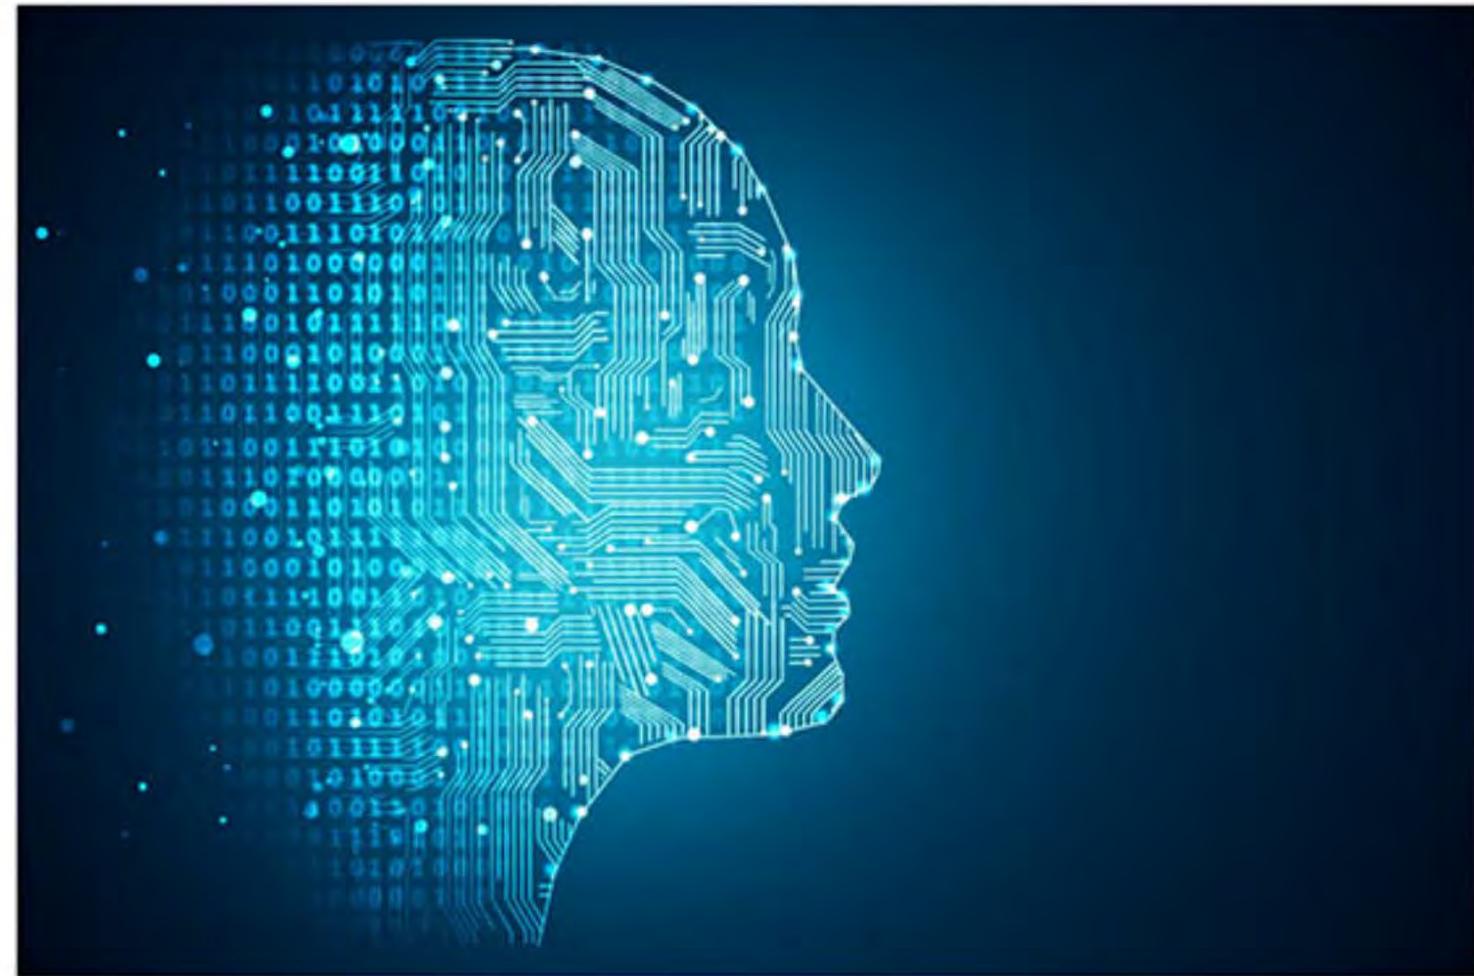


# A Hands-On Introduction to Machine Learning



Julian Gold  
Jonathan Hanke

January 13–16, 2026



*With materials from:*

Brian Arnold, Gage DeZoort, Julian Gold, Jonathan Halverson, Jonathan Hanke, Christina Peters  
Jake Snell, Savannah Thias, Amy Winecoff



# Mini-Course Outline

Date	Topic	Instructor
Tue. 1/13	Machine Learning Overview and Simple Models	Julian Gold
Wed. 1/14	Model Evaluation and Improving Performance	Julian Gold
Thu. 1/15	Introduction to Neural Networks	Jonathan Hanke
Fri. 1/16	Survey of Neural Network Architectures	Jonathan Hanke

## **Artificial Intelligence**

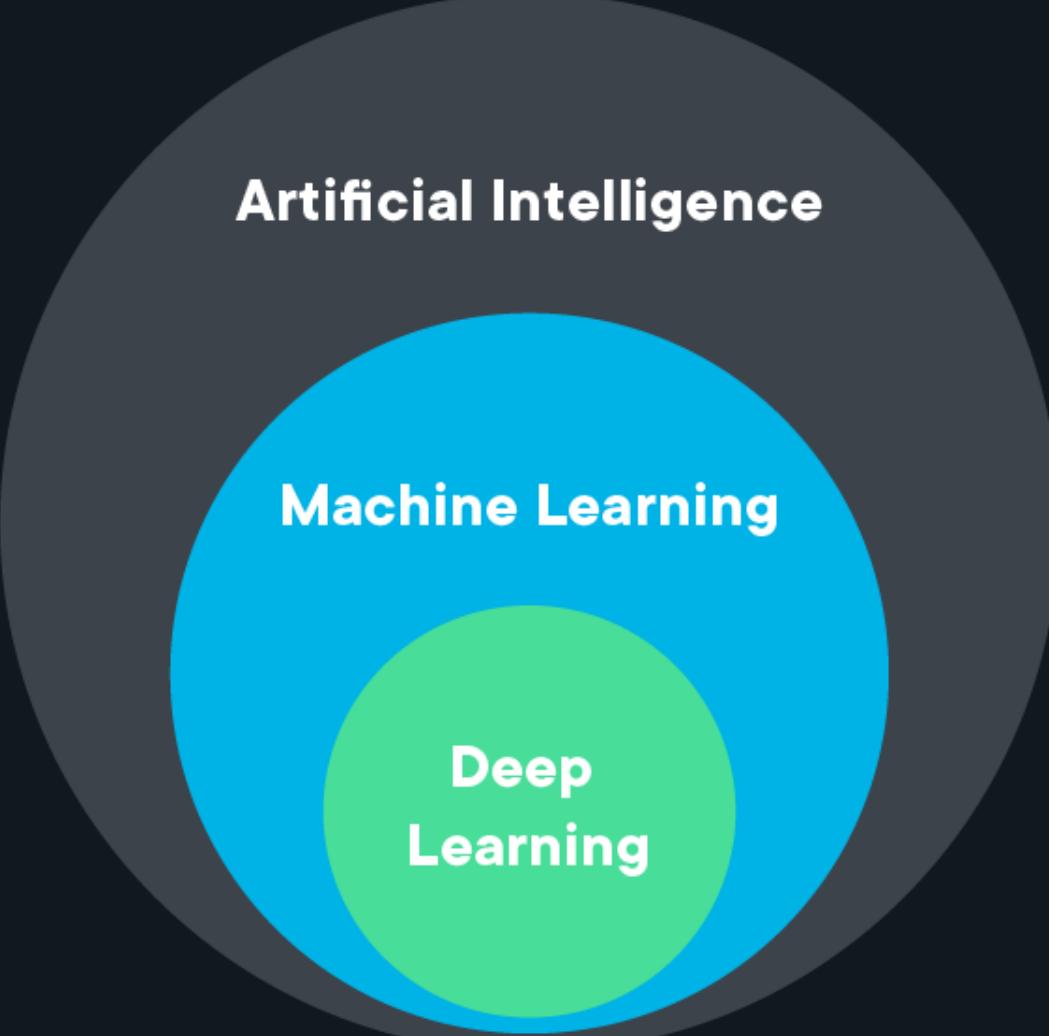
A science devoted to making machines think and act like humans.

## **Machine Learning**

Focuses on enabling computers to perform tasks without explicit programming.

## **Deep Learning**

A subset of machine learning based on artificial neural networks.



**Artificial Intelligence**

**Machine Learning**

**Deep  
Learning**

//

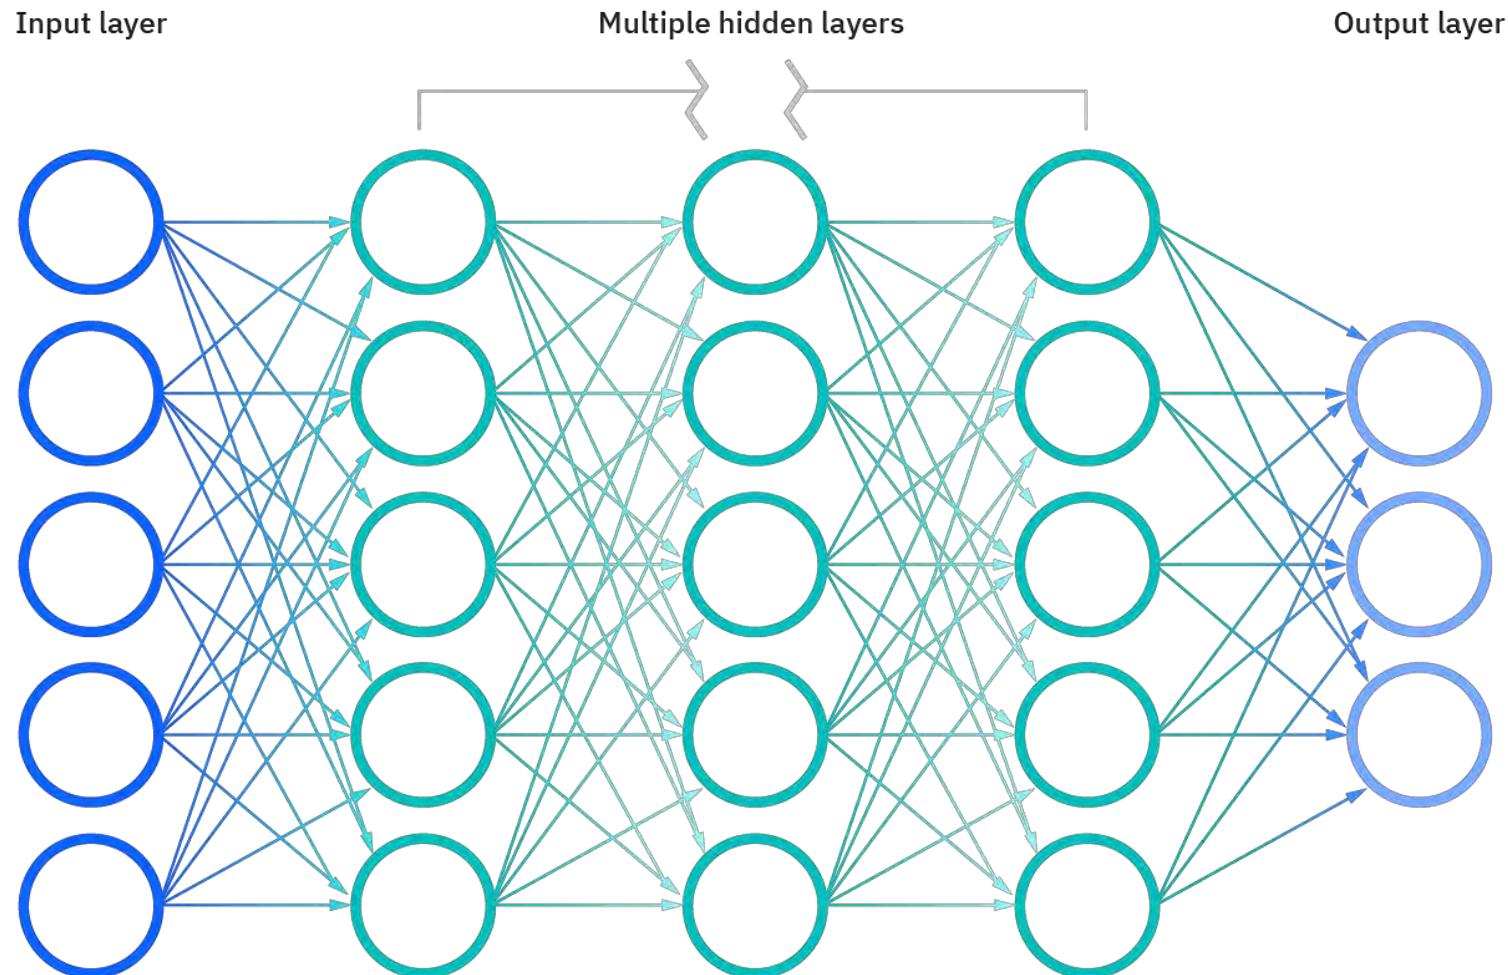
# WHY USE DL?

- Data with complex (highly non-linear) relationships
- Big data – many examples to leverage
- High-dimensional data
- Data with complicated structure (images, video, language, social networks etc.)

## 20 DEEP LEARNING Applications

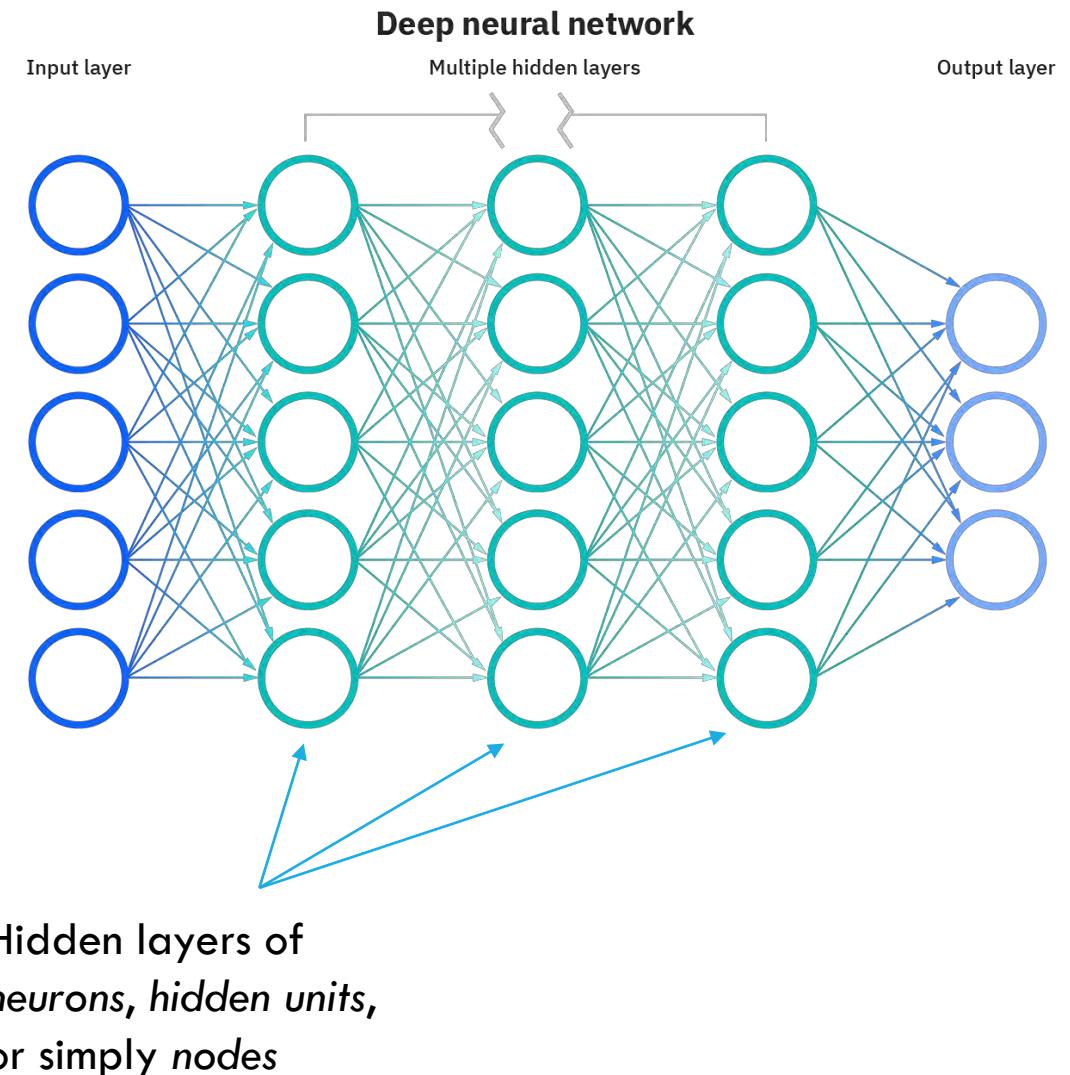


## Deep neural network



# DEEP NEURAL NETWORKS

- A class of ML algorithms based on *artificial neural networks* (ANNs)
  - Neural network → networks of neurons responding to stimuli
  - “Deep” → multiple layers of neurons interacting in sequence

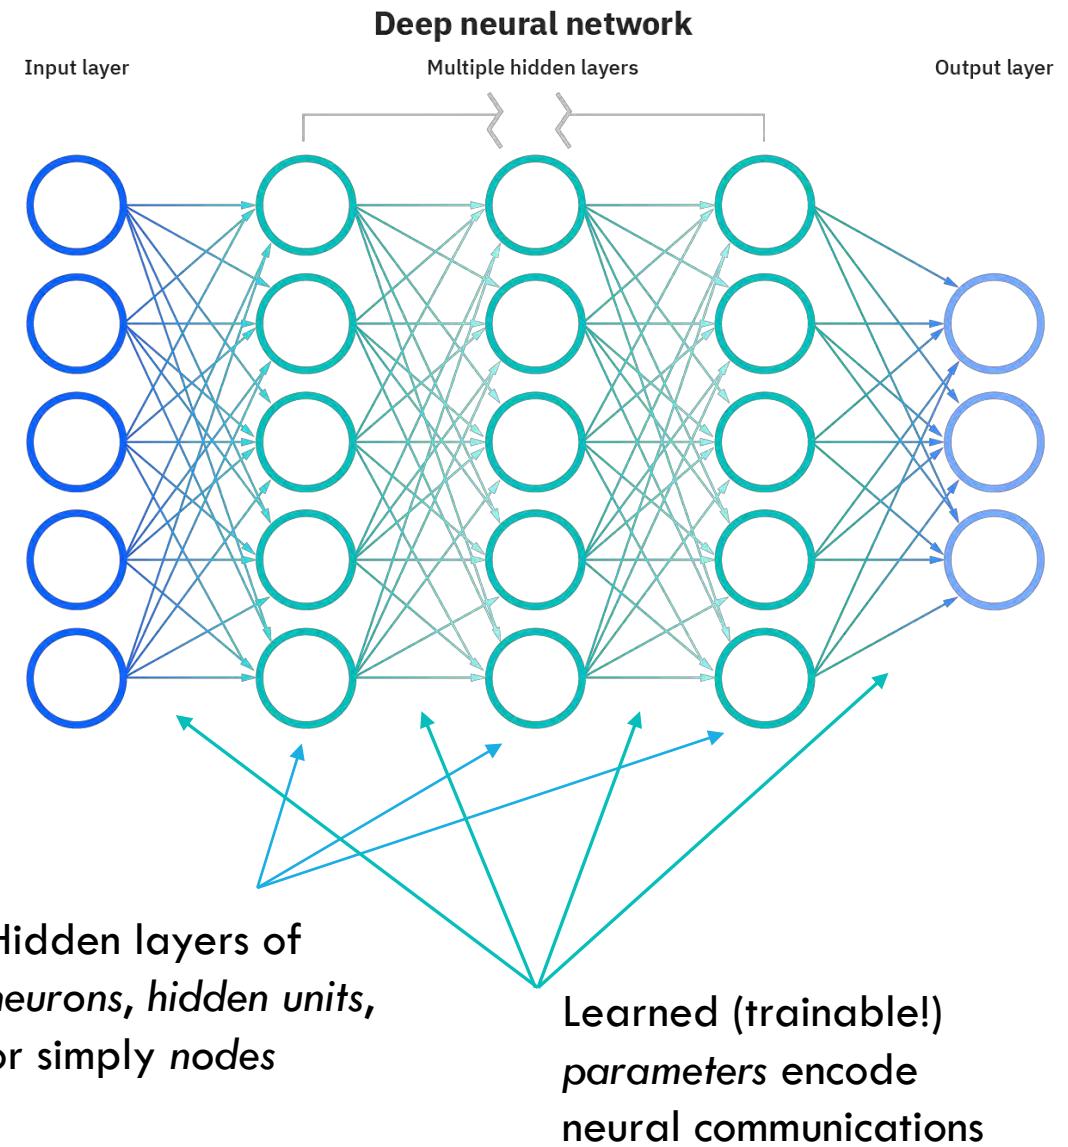


# DEEP NEURAL NETWORKS

- A class of ML algorithms based on *artificial neural networks* (ANNs)
  - Neural network → networks of neurons responding to stimuli
  - “Deep” → multiple layers of neurons interacting in sequence
- Practically speaking, DNNs are non-linear models designed to leverage complicated relationships in data

$$f(x \mid \text{parameters}) = \text{output}$$

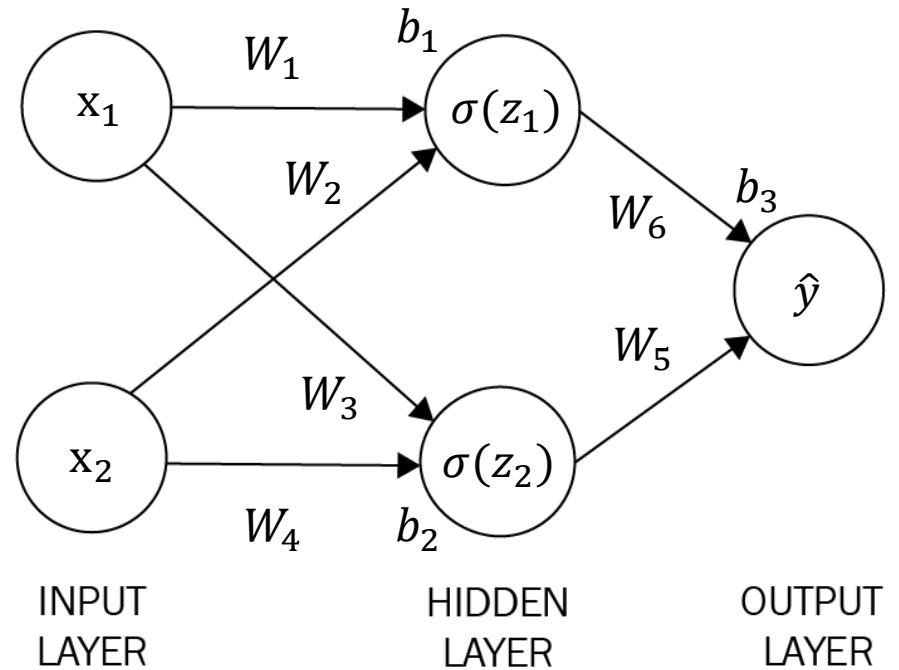
Adjustable, e.g. fit to data in supervised learning



# SIMPLE NN

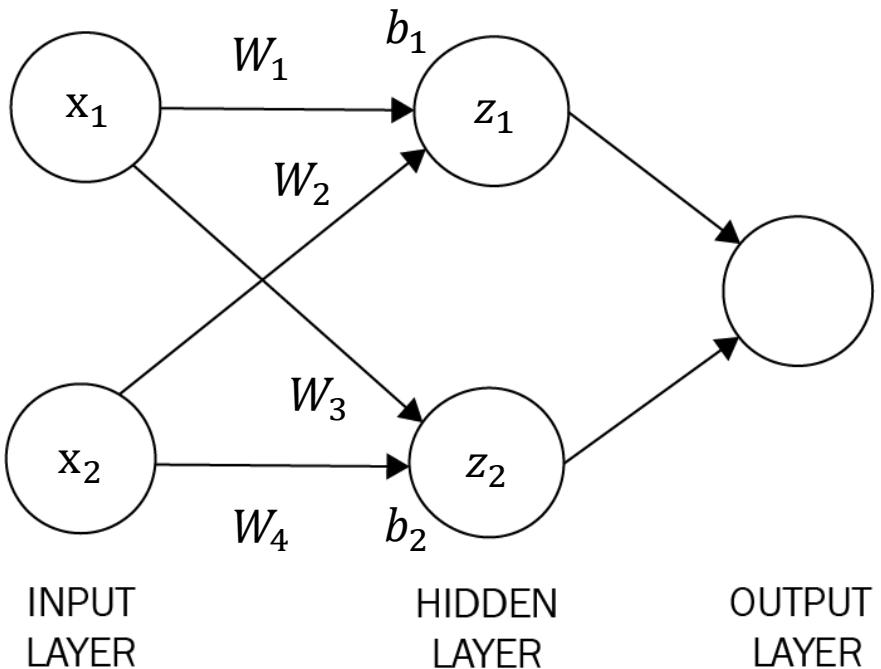
- **Example:** Classification problem with  $x_i \in \mathbb{R}^2$   $y_i \in \{0,1\}$
- Let's draw a test data point and pass it through a simple ANN:

$$x = (x_1, x_2) \quad y = 1$$



- Single hidden layer with 2 neurons
- **Key Ingredients:**
  - Trainable weights  $W_1, W_2, W_3, W_4, W_5, W_6$  and biases  $b_1, b_2, b_3$
  - Non-linear activation functions (called non-linearities)  $\sigma(z)$

# SIMPLE NN



## 1. Compute *preactivations* at each neuron

$$z_1 = w_1 x_1 + w_2 x_2 + b_1$$

$$z_2 = w_3 x_1 + w_4 x_2 + b_2$$

Or, in matrix notation:

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

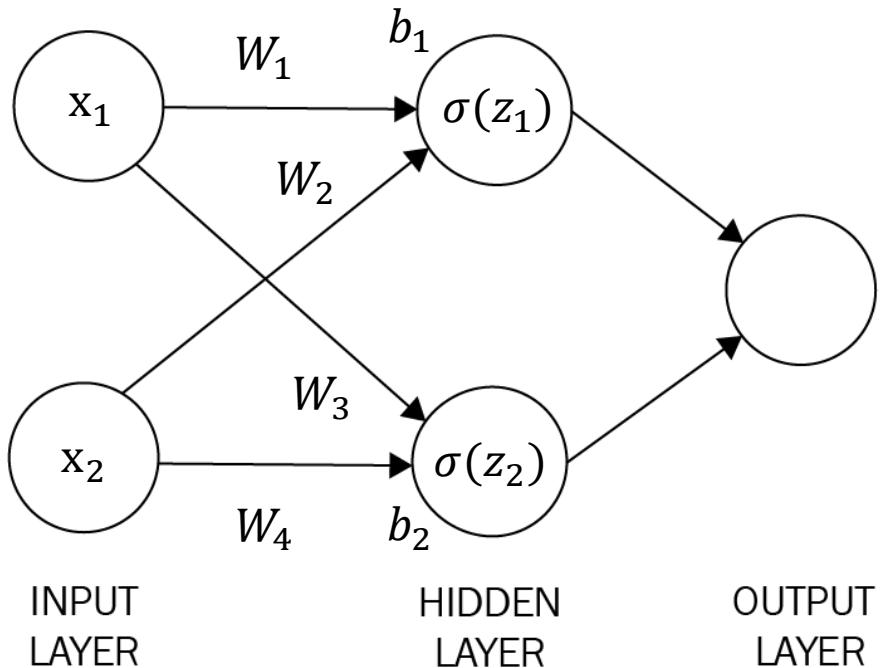


the *weight matrix* mixes up  
the inputs and feeds them  
to the each neuron

the *bias vector* adds  
constants to the mixed  
inputs

**Note: this is a linear operation!!**

# SIMPLE NN



**1. Compute *preactivations* at each neuron**

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

**2. Calculate how much the neuron activates given the strength of the pre-activation**

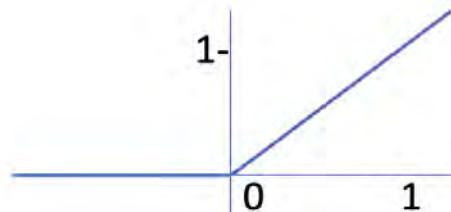
$\sigma(z) \rightarrow$  activation function

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$

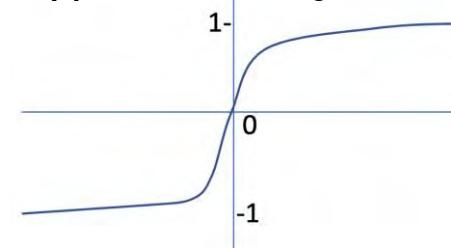
# NON-LINEAR ACTIVATION FUNCTIONS

- Non-linear activation functions allow us to learn complex relationships by “intervening” between linear operations
- In practice, they allow neurons to switch “on” and “off” to varying degrees

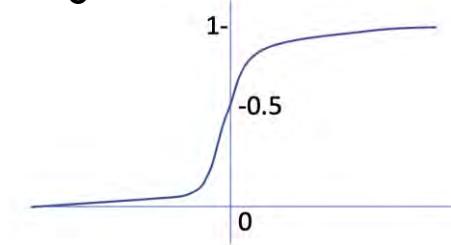
Rectified Linear Unit (ReLU)



Hyperbolic Tangent

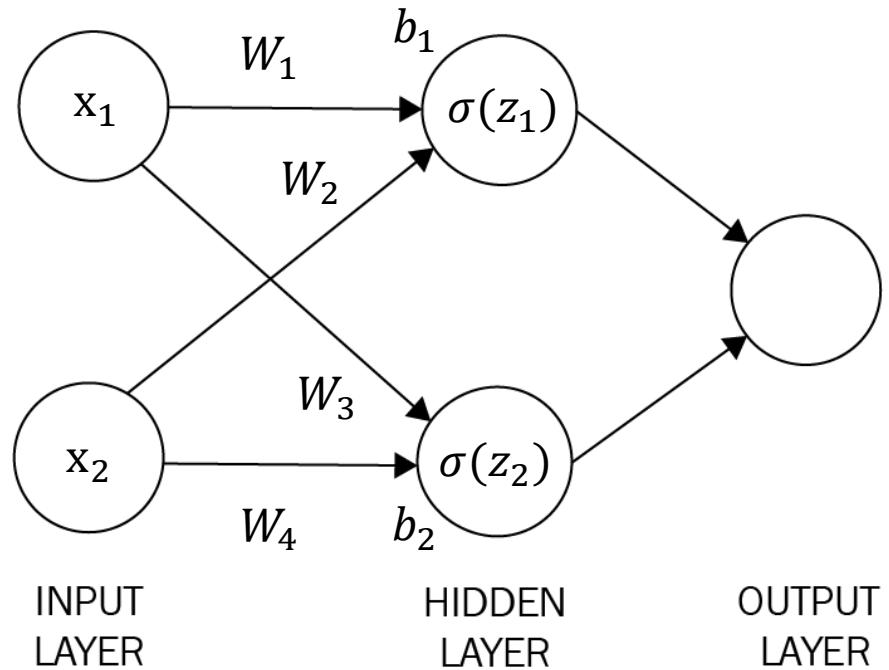


Sigmoid



← Most popular choice; simple to compute (fast training), no “saturating” regions with tiny gradients

# SIMPLE NN



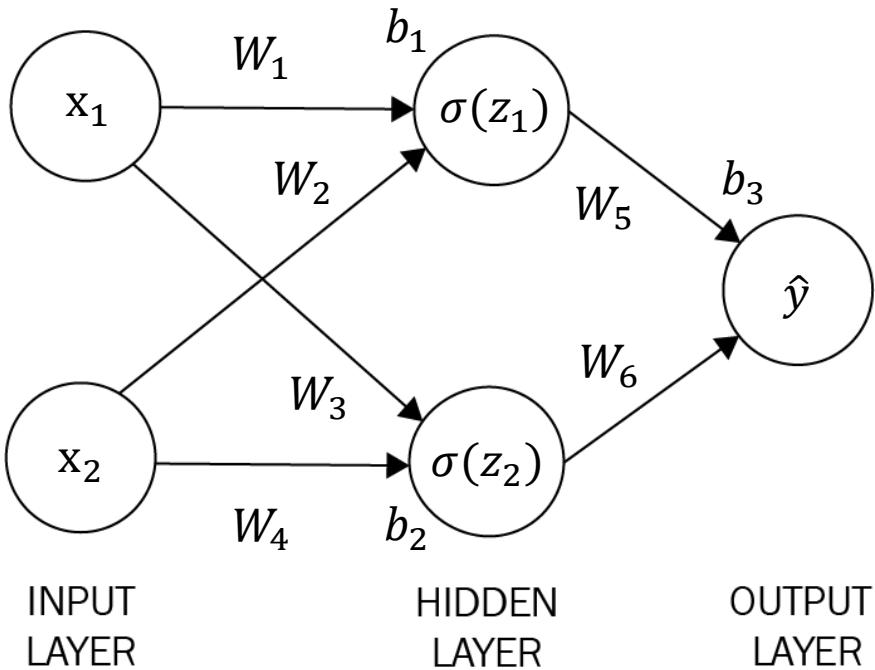
**1. Compute preactivations at each neuron**

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

**2. Calculate how much the neuron activates given the strength of the preactivation**

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$

# HOW IT WORKS



**1. Compute preactivations at each neuron**

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} w_1 & w_2 \\ w_3 & w_4 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

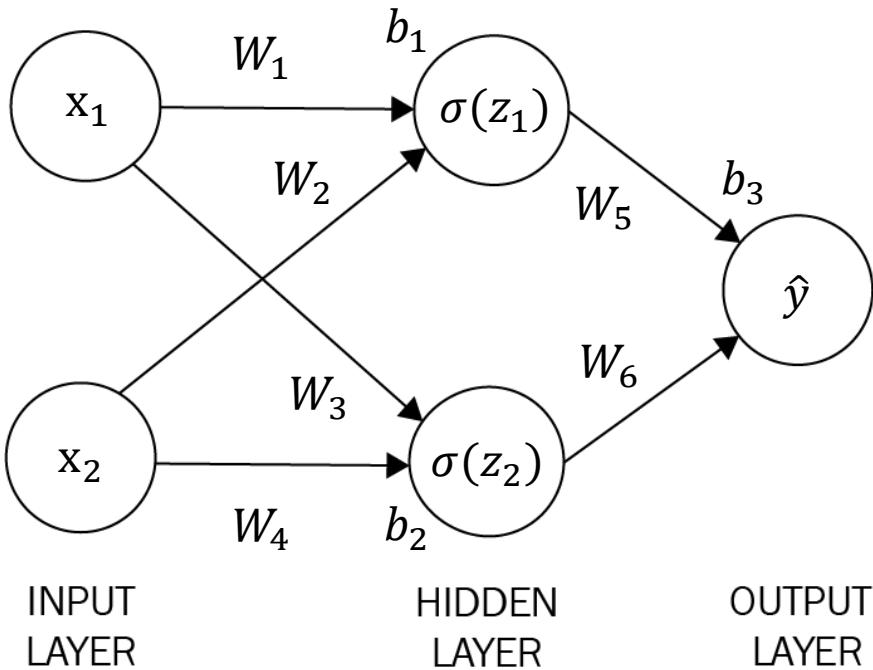
**2. Calculate how much the neuron activates given the strength of the preactivation**

$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} \rightarrow \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix}$$

**3. Calculate model outputs**

$$\hat{y} = (w_5 \quad w_6) \begin{pmatrix} \sigma(z_1) \\ \sigma(z_2) \end{pmatrix} + b_3$$

# HOW IT WORKS



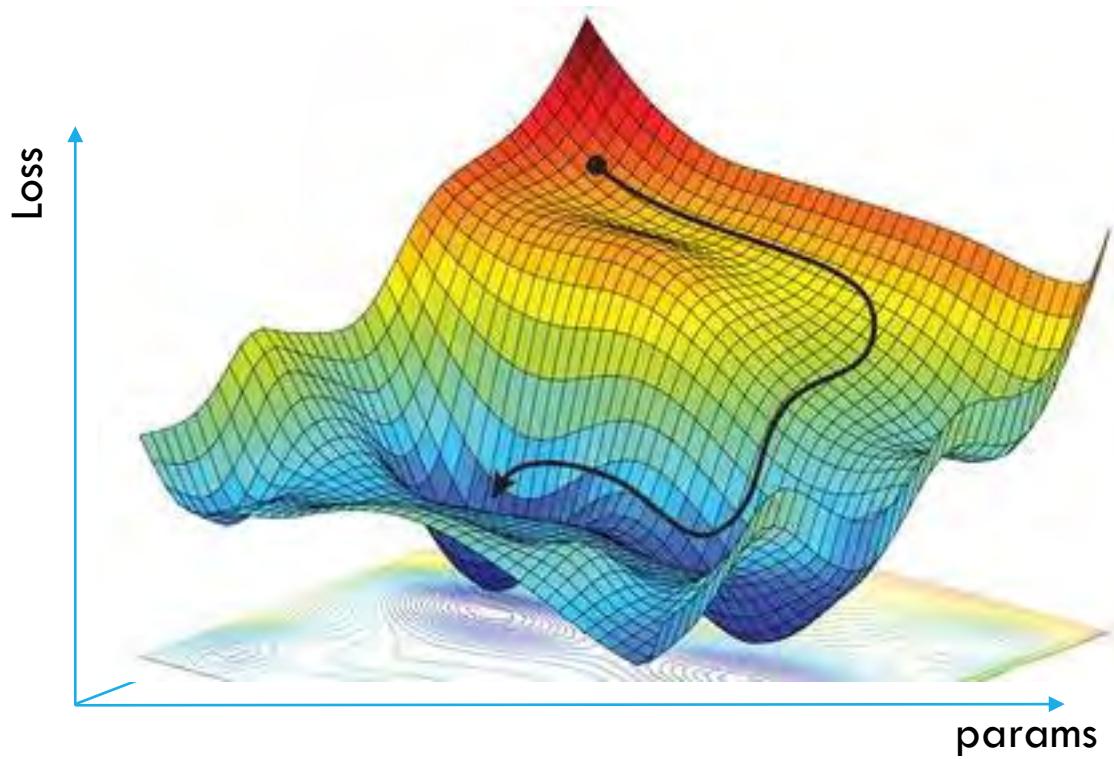
## Forward Pass → Predictions

1. Compute *preactivations* at each neuron
2. Calculate how much the neuron activates given the strength of the *pre-activation*  
*(Repeat 1 and 2 to some fixed depth)*
3. Calculate model outputs

When we define an NN, we fix an architecture by specifying:

- Number of hidden layers (here 1)
- Dimension of the hidden layers (here 2)
- Activation functions
- Random initial values for weights and biases

# TRAINING A NN



The loss function may be very complicated in practice!

Training a NN → find “optimal” weights, biases

- Start by defining a **loss function**  $L(y, \hat{y})$
- Compute the gradient of  $L(y, \hat{y})$  with respect to each weight and bias
- Update the weights and biases via **gradient descent**:

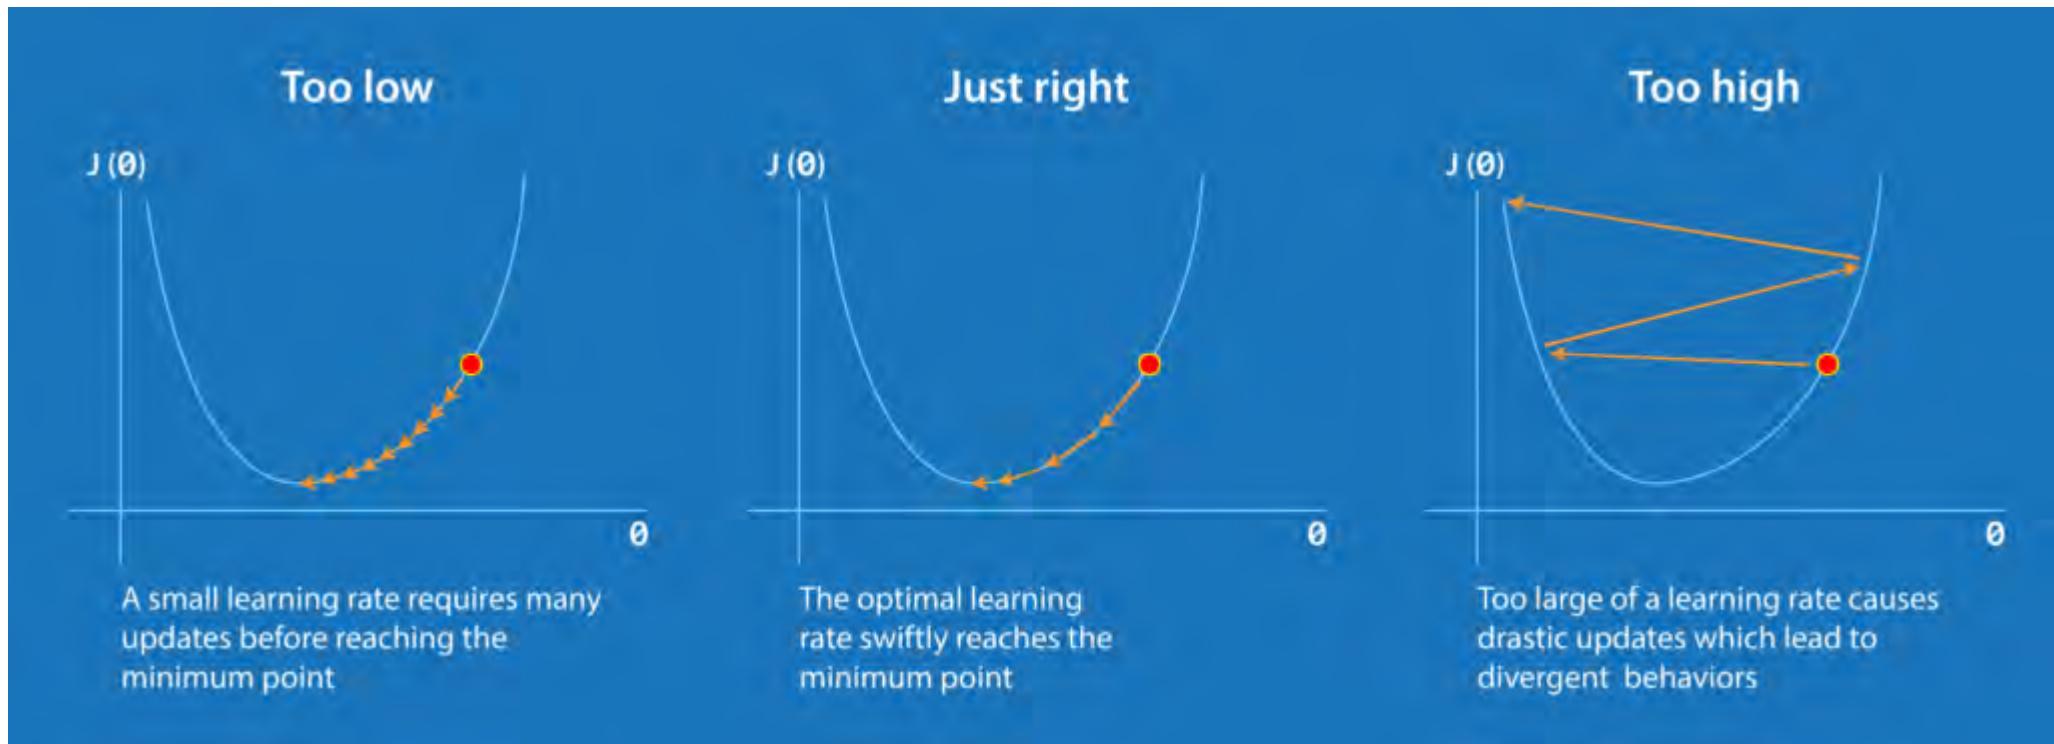
$$W_1^{(k+1)} = W_1^{(k)} - \gamma \frac{\partial L}{\partial W_1} \Big|_{W_1^{(k)}}$$

$$b_1^{(k+1)} = b_1^{(k)} - \gamma \frac{\partial L}{\partial b_1} \Big|_{b_1^{(k)}}$$

etc.

$\gamma \rightarrow$  learning rate

# LEARNING RATES



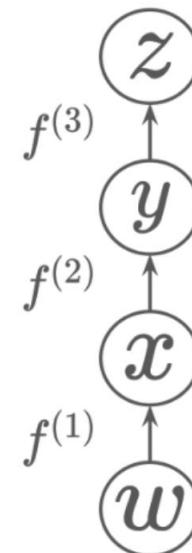
# TRAINING A NN

## Backward Pass (Backpropagation)

Apply the chain rule to calculate derivatives of the loss with respect to the weights/biases

What we want:  $\frac{\partial L}{\partial W}, \frac{\partial L}{\partial b}$

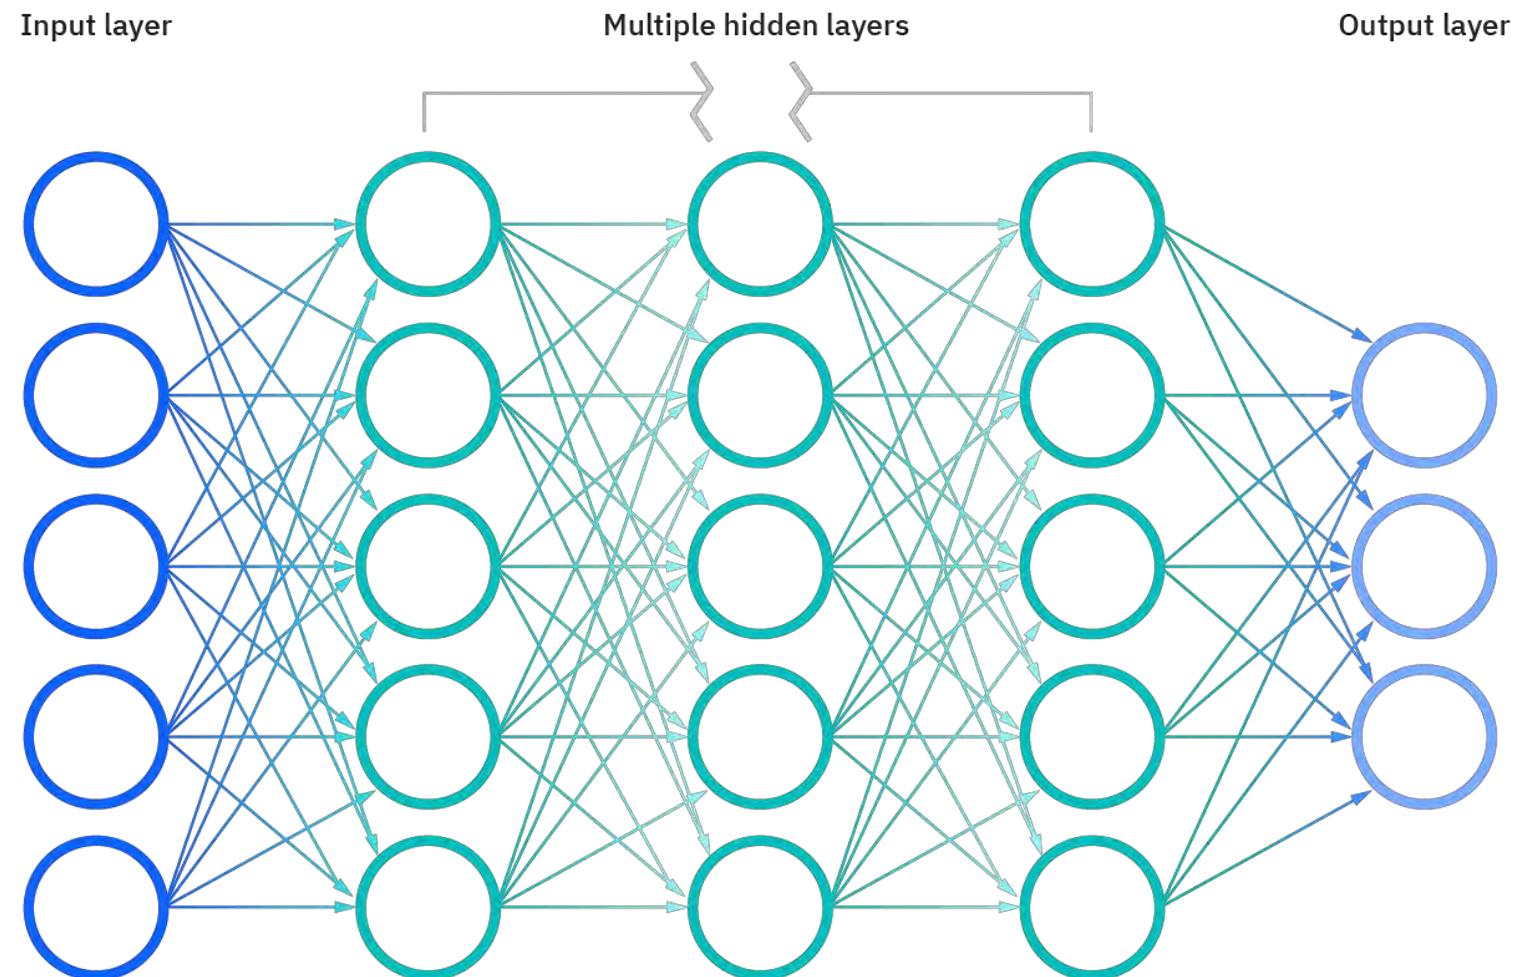
Generically, “gradients”


$$\begin{aligned}\frac{\partial z}{\partial w} &= \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w} \\ &= f^{(3)'}(y)f^{(2)'}(x)f^{(1)'}(w)\end{aligned}$$

# PSEUDOCODE

```
for each training epoch:  
    for each (input, truth) in train_data:  
        prediction = NN(train_data)  
        loss = loss_function(prediction, truth)  
        gradients = compute_gradients(loss)  
        new_params = grad_descent(gradients, NN.parameters, lr)  
        model.update(new_params)
```

## Deep neural network



# SHOULD YOU USE DL FOR YOUR PROJECT?

## Why DL?

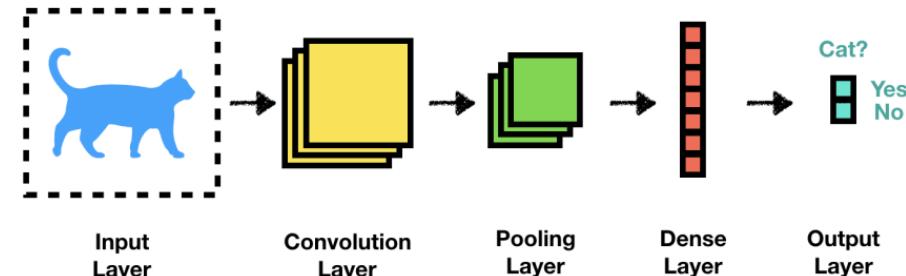
- Produce predictions in multiple feed-forward stages → powerful feature extraction:
  - Less need for data pre-processing
  - Ability to leverage large amounts of data (“big data”)
- Handle more complicated data representations like images, sentences, and graphs
- Can be designed to perform complicated (multi-stage) tasks end-to-end

## Why not DL?

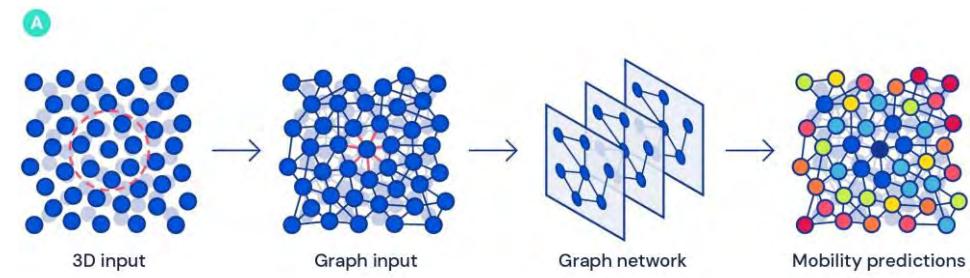
- Developing and training large DL models is computationally expensive (slow, resource intensive) and often requires specialized computing hardware
- It usually takes a lot of data to train DL models

# MORE ARCHITECTURES

- NNs are the building blocks for more complicated architectures, e.g.
  - **Convolutional Neural Networks** are frequently applied to images or other grid data
  - **Recurrent Neural Networks** are applied to sequences like sentences
  - **Graph Neural Networks** operate on networks of objects (nodes) connected by their relationships (edges)
  - **Generative Adversarial Networks** are used to generate new data (e.g. photographs) similar to a reference set



[Convolutional Neural Network: A Step By Step Guide | by Shashikant | Towards Data Science](#)



[Towards understanding glasses with graph neural networks \(deepmind.com\)](#)

A mostly complete chart of

# Neural Networks

©2016 Fjodor van Veer - asimovinstitute.org

- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

Deep Feed Forward (DFF)



Perceptron (P)

Feed Forward (FF)

Radial Basis Network (RBF)

Recurrent Neural Network (RNN)

Long / Short Term Memory (LSTM)

Gated Recurrent Unit (GRU)

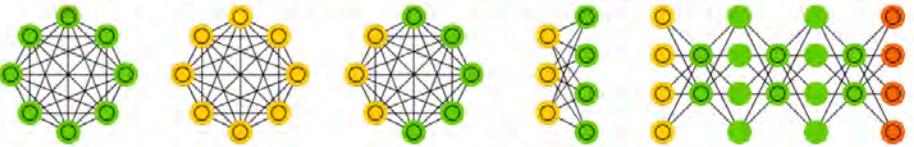
Auto Encoder (AE)

Variational AE (VAE)

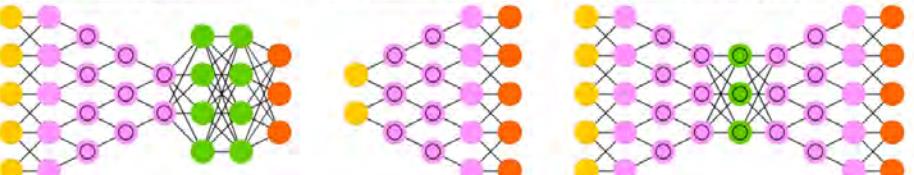
Denoising AE (DAE)

Sparse AE (SAE)

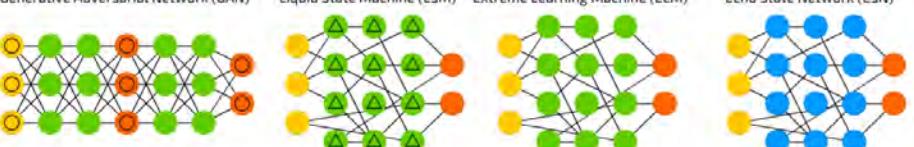
Markov Chain (MC) Hopfield Network (HN) Boltzmann Machine (BM) Restricted BM (RBM) Deep Belief Network (DBN)



Deep Convolutional Network (DCN) Deconvolutional Network (DN) Deep Convolutional Inverse Graphics Network (DCIGN)



Generative Adversarial Network (GAN) Liquid State Machine (LSM) Extreme Learning Machine (ELM) Echo State Network (ESN)

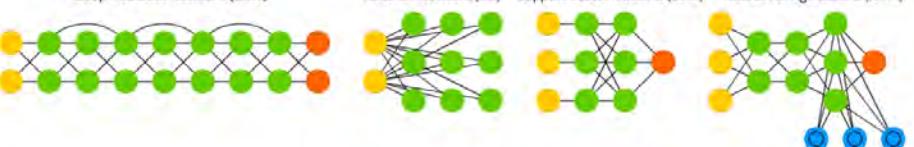


Deep Residual Network (DRN)

Kohonen Network (KN)

Support Vector Machine (SVM)

Neural Turing Machine (NTM)



A bit outdated, but fun to see the creativity...

[The mostly complete chart of Neural Networks, explained | by Andrew Tch | Towards Data Science](#)



**Time for some NN practice!**

Please navigate to Day 3!

[PrincetonUniversity/intro\\_machine\\_learning \(github.com\)](https://github.com/PrincetonUniversity/intro_machine_learning)