



CHECK-IN

INTRODUCTION TO MACHINE LEARNING (PART 4 OF 5)

MONDAY, JANUARY 22 AT 2:00PM



<https://cglink.me/2gi/c19387711116154454>

- ① Open the **My PrincetonU** app.
- ② Select a Hub
- ③ Click on QR Code scanner.
- ④ Scan this QR Code and you are checked-in!



Day 5 Hackathon

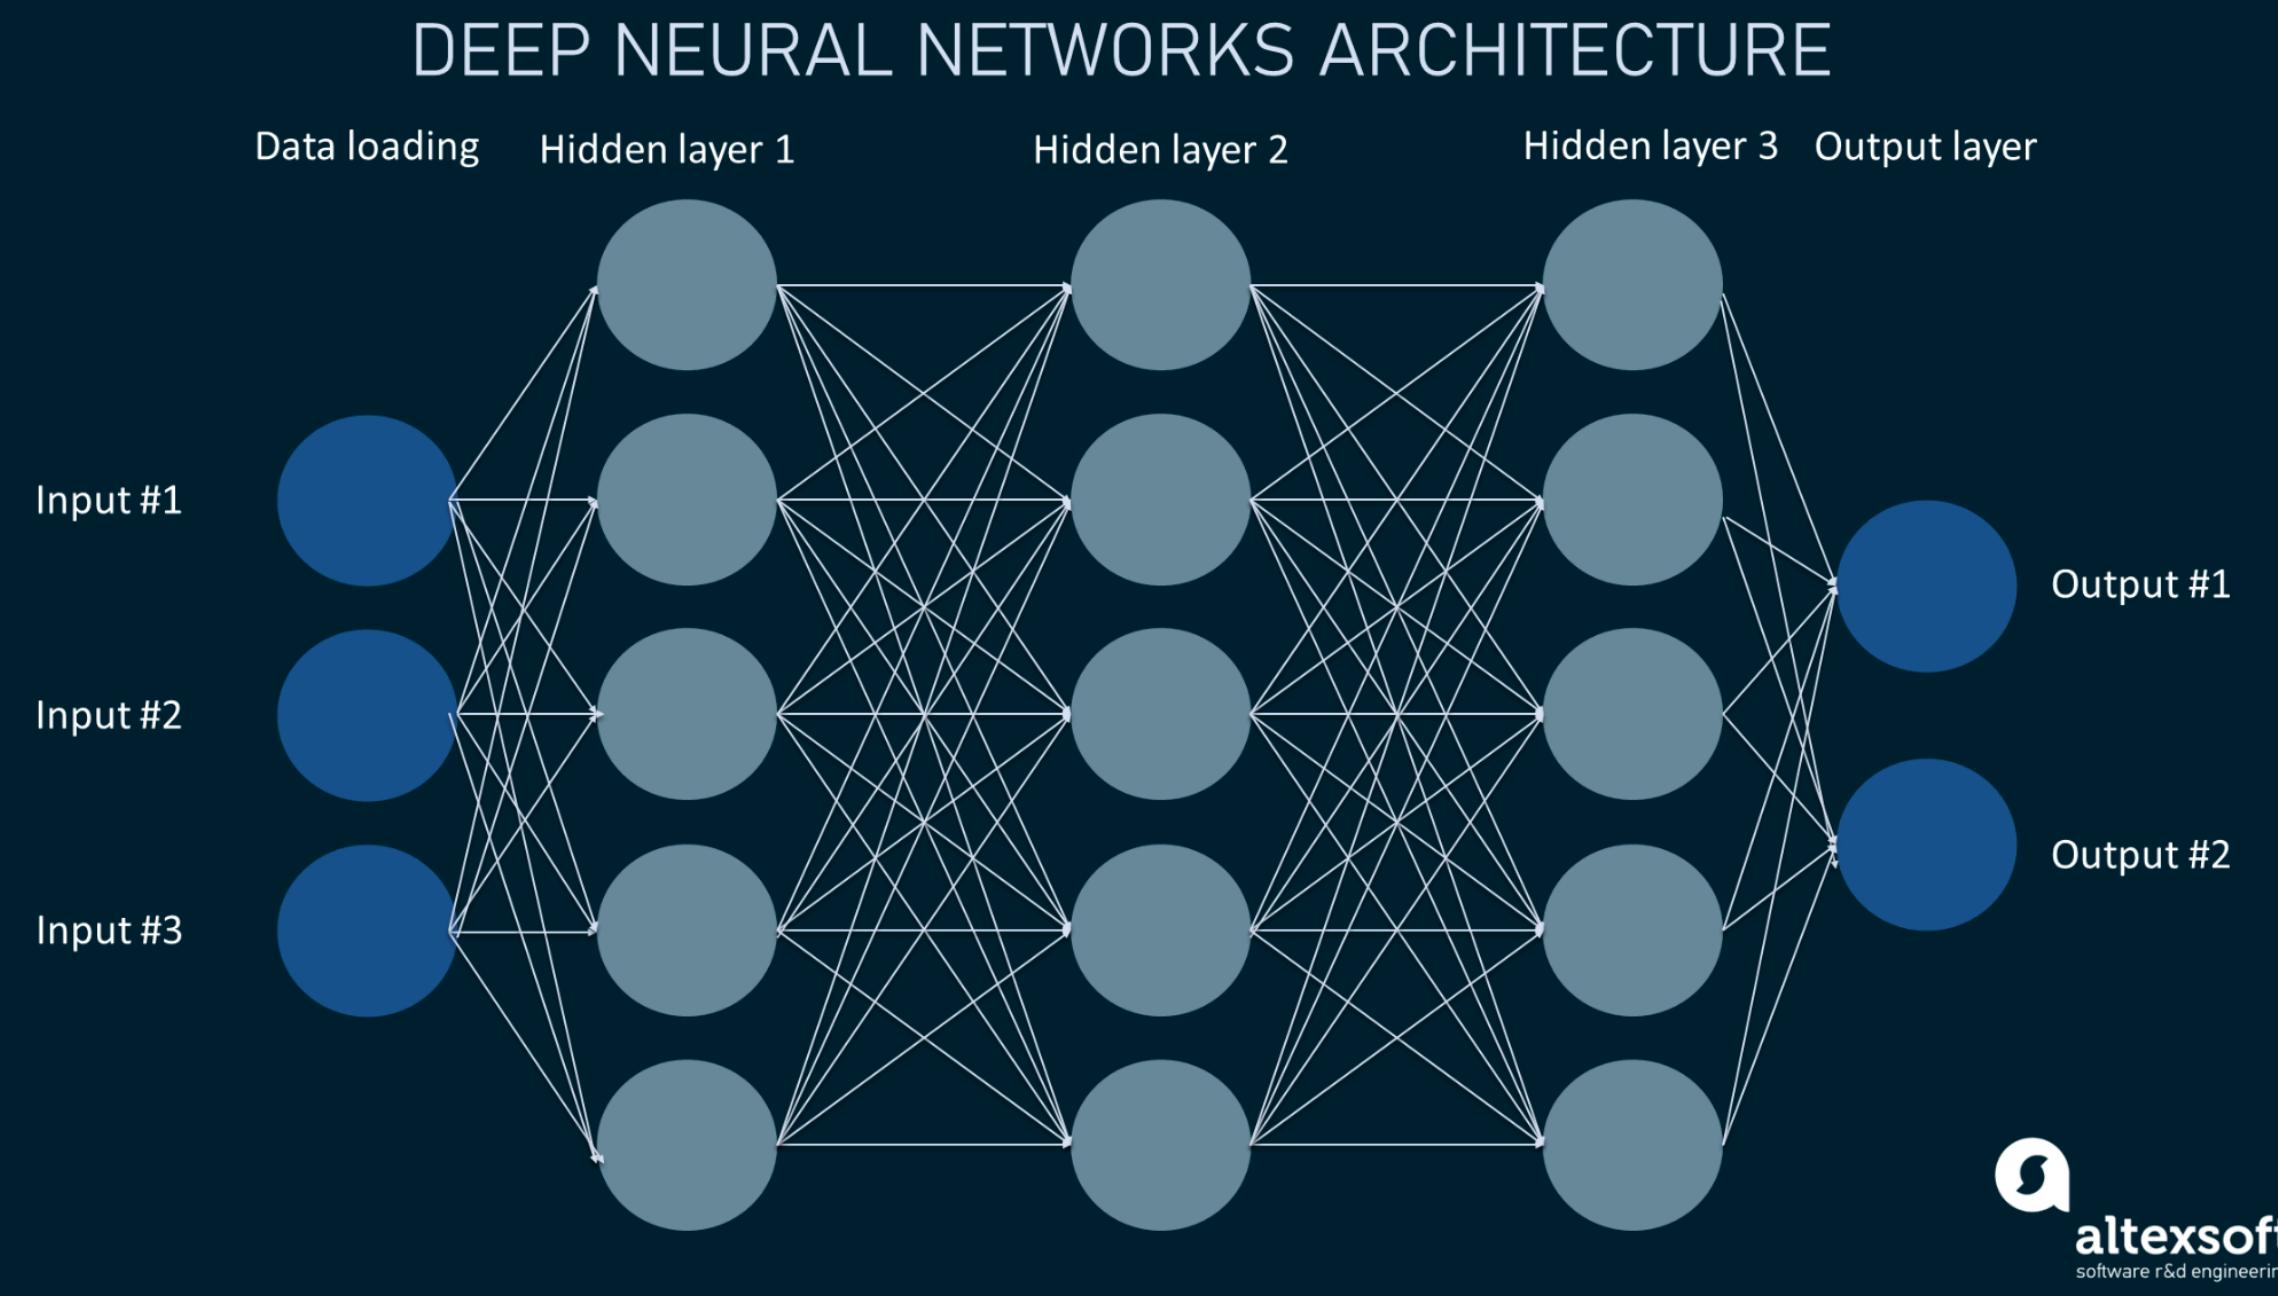
Project Descriptions

- **Computer Vision:** Learn more about CNNs, classify dogs versus cats using a simple CNN, and use transfer learning with an advanced CNN (ResNet-50) to classify dogs versus cats.
- **Diffusion Models:** Learn about diffusion models (e.g., DALL-E 2) then build one and train a generative model for images.
- **Large Language Models:** This session introduces the basics of language modeling using the transformer architecture. Participants will learn how to download and fine-tune an LLM using the Hugging Face library.

Deep Learning Architectures

Introduction to Machine Learning, Day 4

DEEP NEURAL NETWORKS ARCHITECTURE



- Forward pass:

$$z_i^{(\ell+1)} = \sum_{j=1}^{n_\ell} W_{ij}^{(\ell+1)} \sigma(z_j^{(\ell)}) + b_i^{(\ell+1)}$$

- Backward pass:

$$W_{ij}^{(\ell+1)} = W_{ij}^{(\ell)} - \gamma \frac{\partial L}{\partial W_{ij}} \Big|_{W_{ij}^{(\ell)}}$$

$$b_i^{(\ell+1)} = b_i^{(\ell)} - \gamma \frac{\partial L}{\partial b_i} \Big|_{b_i^{(\ell)}}$$

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

```

def train_one_epoch(epoch_index, tb_writer):
    running_loss = 0.
    last_loss = 0.

    # Here, we use enumerate(training_loader) instead of
    # iter(training_loader) so that we can track the batch
    # index and do some intra-epoch reporting
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        # Gather data and report
        running_loss += loss.item()
        if i % 1000 == 999:
            last_loss = running_loss / 1000 # loss per batch
            print(' batch {} loss: {}'.format(i + 1, last_loss))
            tb_x = epoch_index * len(training_loader) + i + 1
            tb_writer.add_scalar('Loss/train', last_loss, tb_x)
            running_loss = 0.

    return last_loss

```

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

$$z_i^{(\ell+1)} = \sum_{j=1}^{n_\ell} W_{ij}^{(\ell+1)} \sigma(z_j^{(\ell)}) + b_i^{(\ell+1)}$$

```

def train_one_epoch(epoch_index, tb_writer):
    running_loss = 0.
    last_loss = 0.

    # Here, we use enumerate(training_loader) instead of
    # iter(training_loader) so that we can track the batch
    # index and do some intra-epoch reporting
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        # Gather data and report
        running_loss += loss.item()
        if i % 1000 == 999:
            last_loss = running_loss / 1000 # loss per batch
            print(' batch {} loss: {}'.format(i + 1, last_loss))
            tb_x = epoch_index * len(training_loader) + i + 1
            tb_writer.add_scalar('Loss/train', last_loss, tb_x)
            running_loss = 0.

    return last_loss

```

```

class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(28*28, 512),
            nn.ReLU(),
            nn.Linear(512, 512),
            nn.ReLU(),
            nn.Linear(512, 10),
        )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

```

$$W_{ij}^{(\ell+1)} = W_{ij}^{(\ell)} - \gamma \frac{\partial L}{\partial W_{ij}} \Big|_{W_{ij}^{(\ell)}} \quad b_i^{(\ell+1)} = b_i^{(\ell)} - \gamma \frac{\partial L}{\partial b_i} \Big|_{b_i^{(\ell)}}$$

```

def train_one_epoch(epoch_index, tb_writer):
    running_loss = 0.
    last_loss = 0.

    # Here, we use enumerate(training_loader) instead of
    # iter(training_loader) so that we can track the batch
    # index and do some intra-epoch reporting
    for i, data in enumerate(training_loader):
        # Every data instance is an input + label pair
        inputs, labels = data

        # Zero your gradients for every batch!
        optimizer.zero_grad()

        # Make predictions for this batch
        outputs = model(inputs)

        # Compute the loss and its gradients
        loss = loss_fn(outputs, labels)
        loss.backward()

        # Adjust learning weights
        optimizer.step()

        # Gather data and report
        running_loss += loss.item()
        if i % 1000 == 999:
            last_loss = running_loss / 1000 # loss per batch
            print(' batch {} loss: {}'.format(i + 1, last_loss))
            tb_x = epoch_index * len(training_loader) + i + 1
            tb_writer.add_scalar('Loss/train', last_loss, tb_x)
            running_loss = 0.

    return last_loss

```

Beyond Simple DNNs

Survey of Deep Learning Architectures

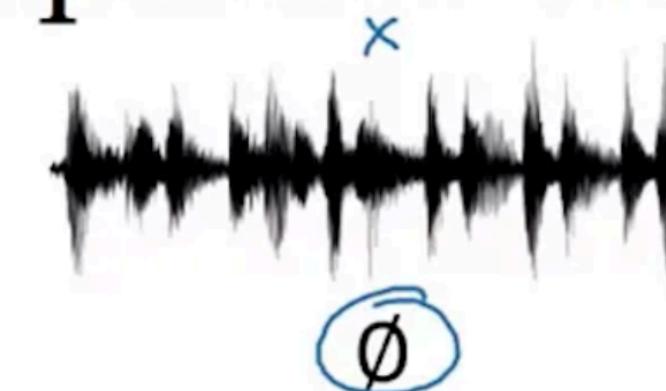
- Deep NN (DNN) \leftrightarrow Feed-Forward NN (FFNN) \leftrightarrow Fully-Connected NN (FCNN)
- Many other architectures exist:
 - Recurrent NNs (RNNs): process sequential data
 - Convolutional NNs (CNNs): process data on a grid
 - Graph Neural Networks (GNNs): process data on a graph / attention
 - Generative Models: produce new data
 - ... and more!

Recurrent Neural Networks

Designed to Process Sequential Data

Examples of sequence data

Speech recognition



"The quick brown fox jumped
over the lazy dog."

Music generation



Sentiment classification

"There is nothing to like
in this movie."



DNA sequence analysis

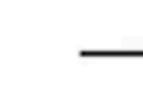
AGCCCCCTGTGAGGAACCTAG



AGCCC_{red}CTGTGAGGAACCTAG

Machine translation

Voulez-vous chanter avec
moi?



Do you want to sing with
me?

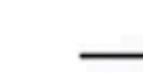
Video activity recognition



Running

Name entity recognition

Yesterday, Harry Potter
met Hermione Granger.

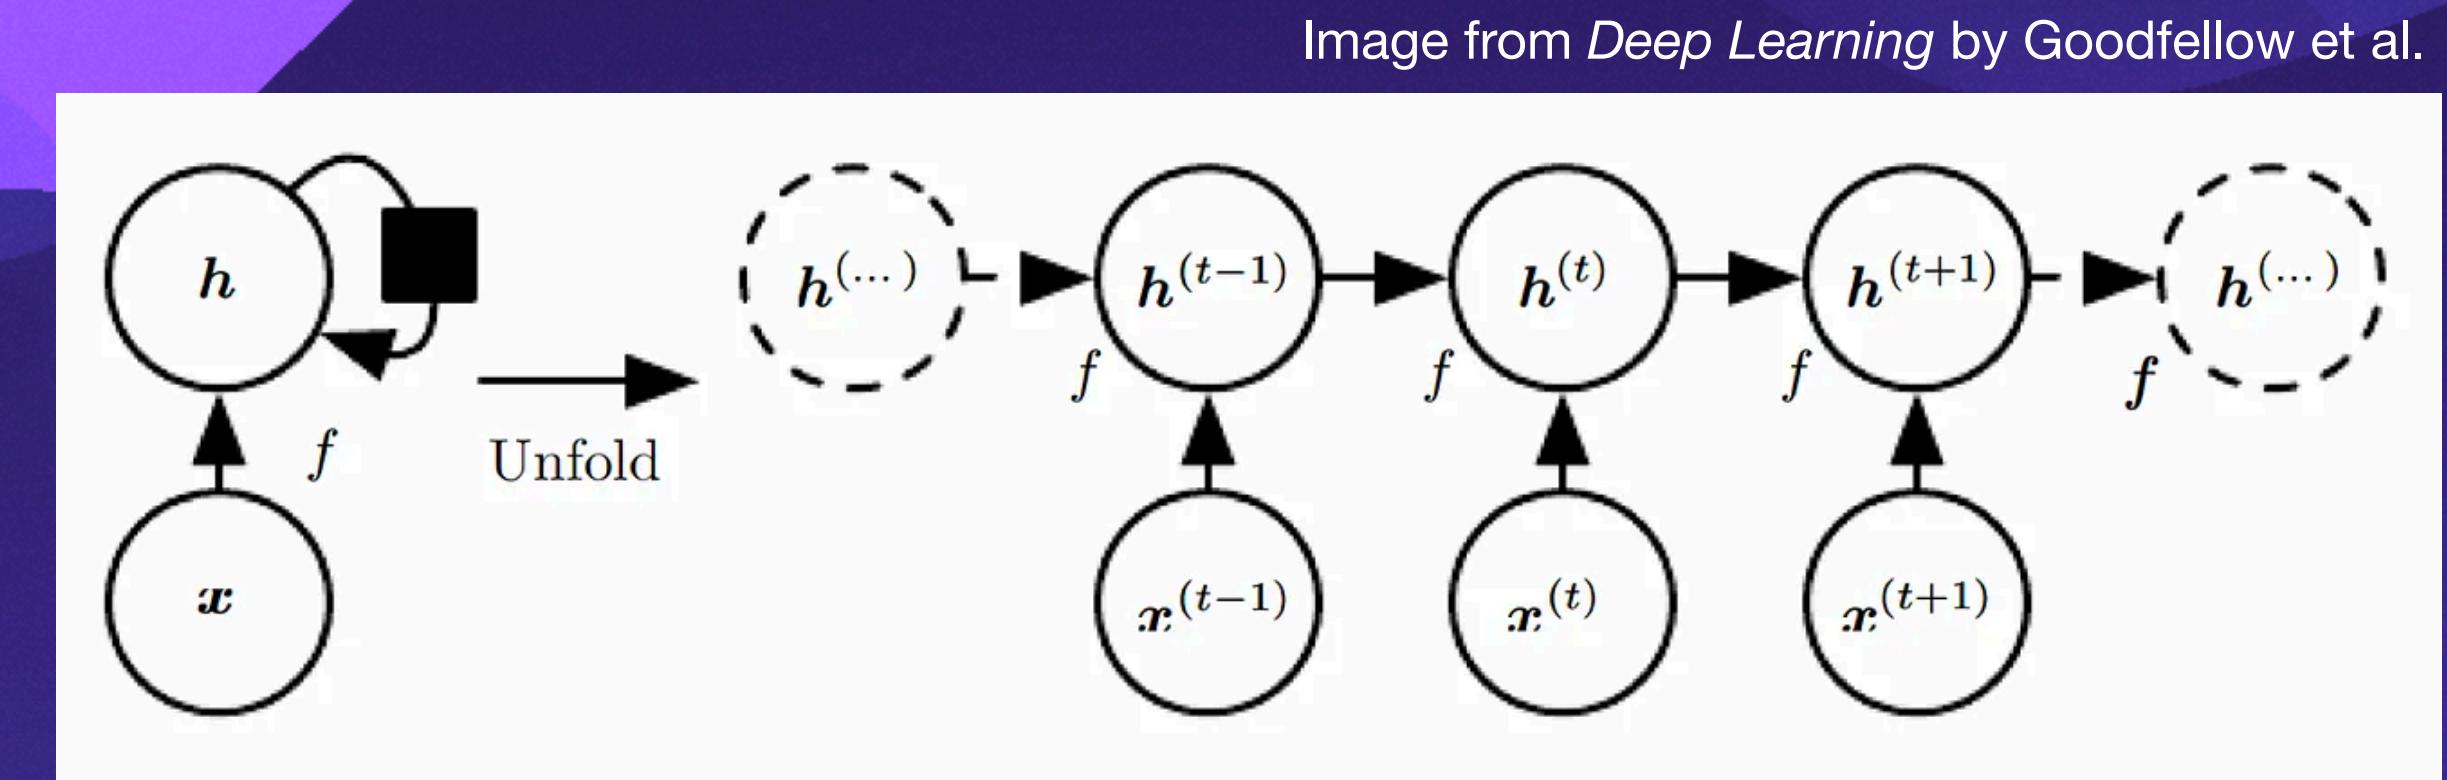


Yesterday, Harry Potter
met Hermione Granger.
Andrew Ng

Recurrent Neural Networks

Basic Idea

- Time-indexed inputs:
 $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(T)}$
- Parameter sharing: apply the same set of learnable weights to all values of the time index
- Given a set of learnable parameters θ , the hidden units in many RNNs are calculated via.
$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)}; \theta)$$



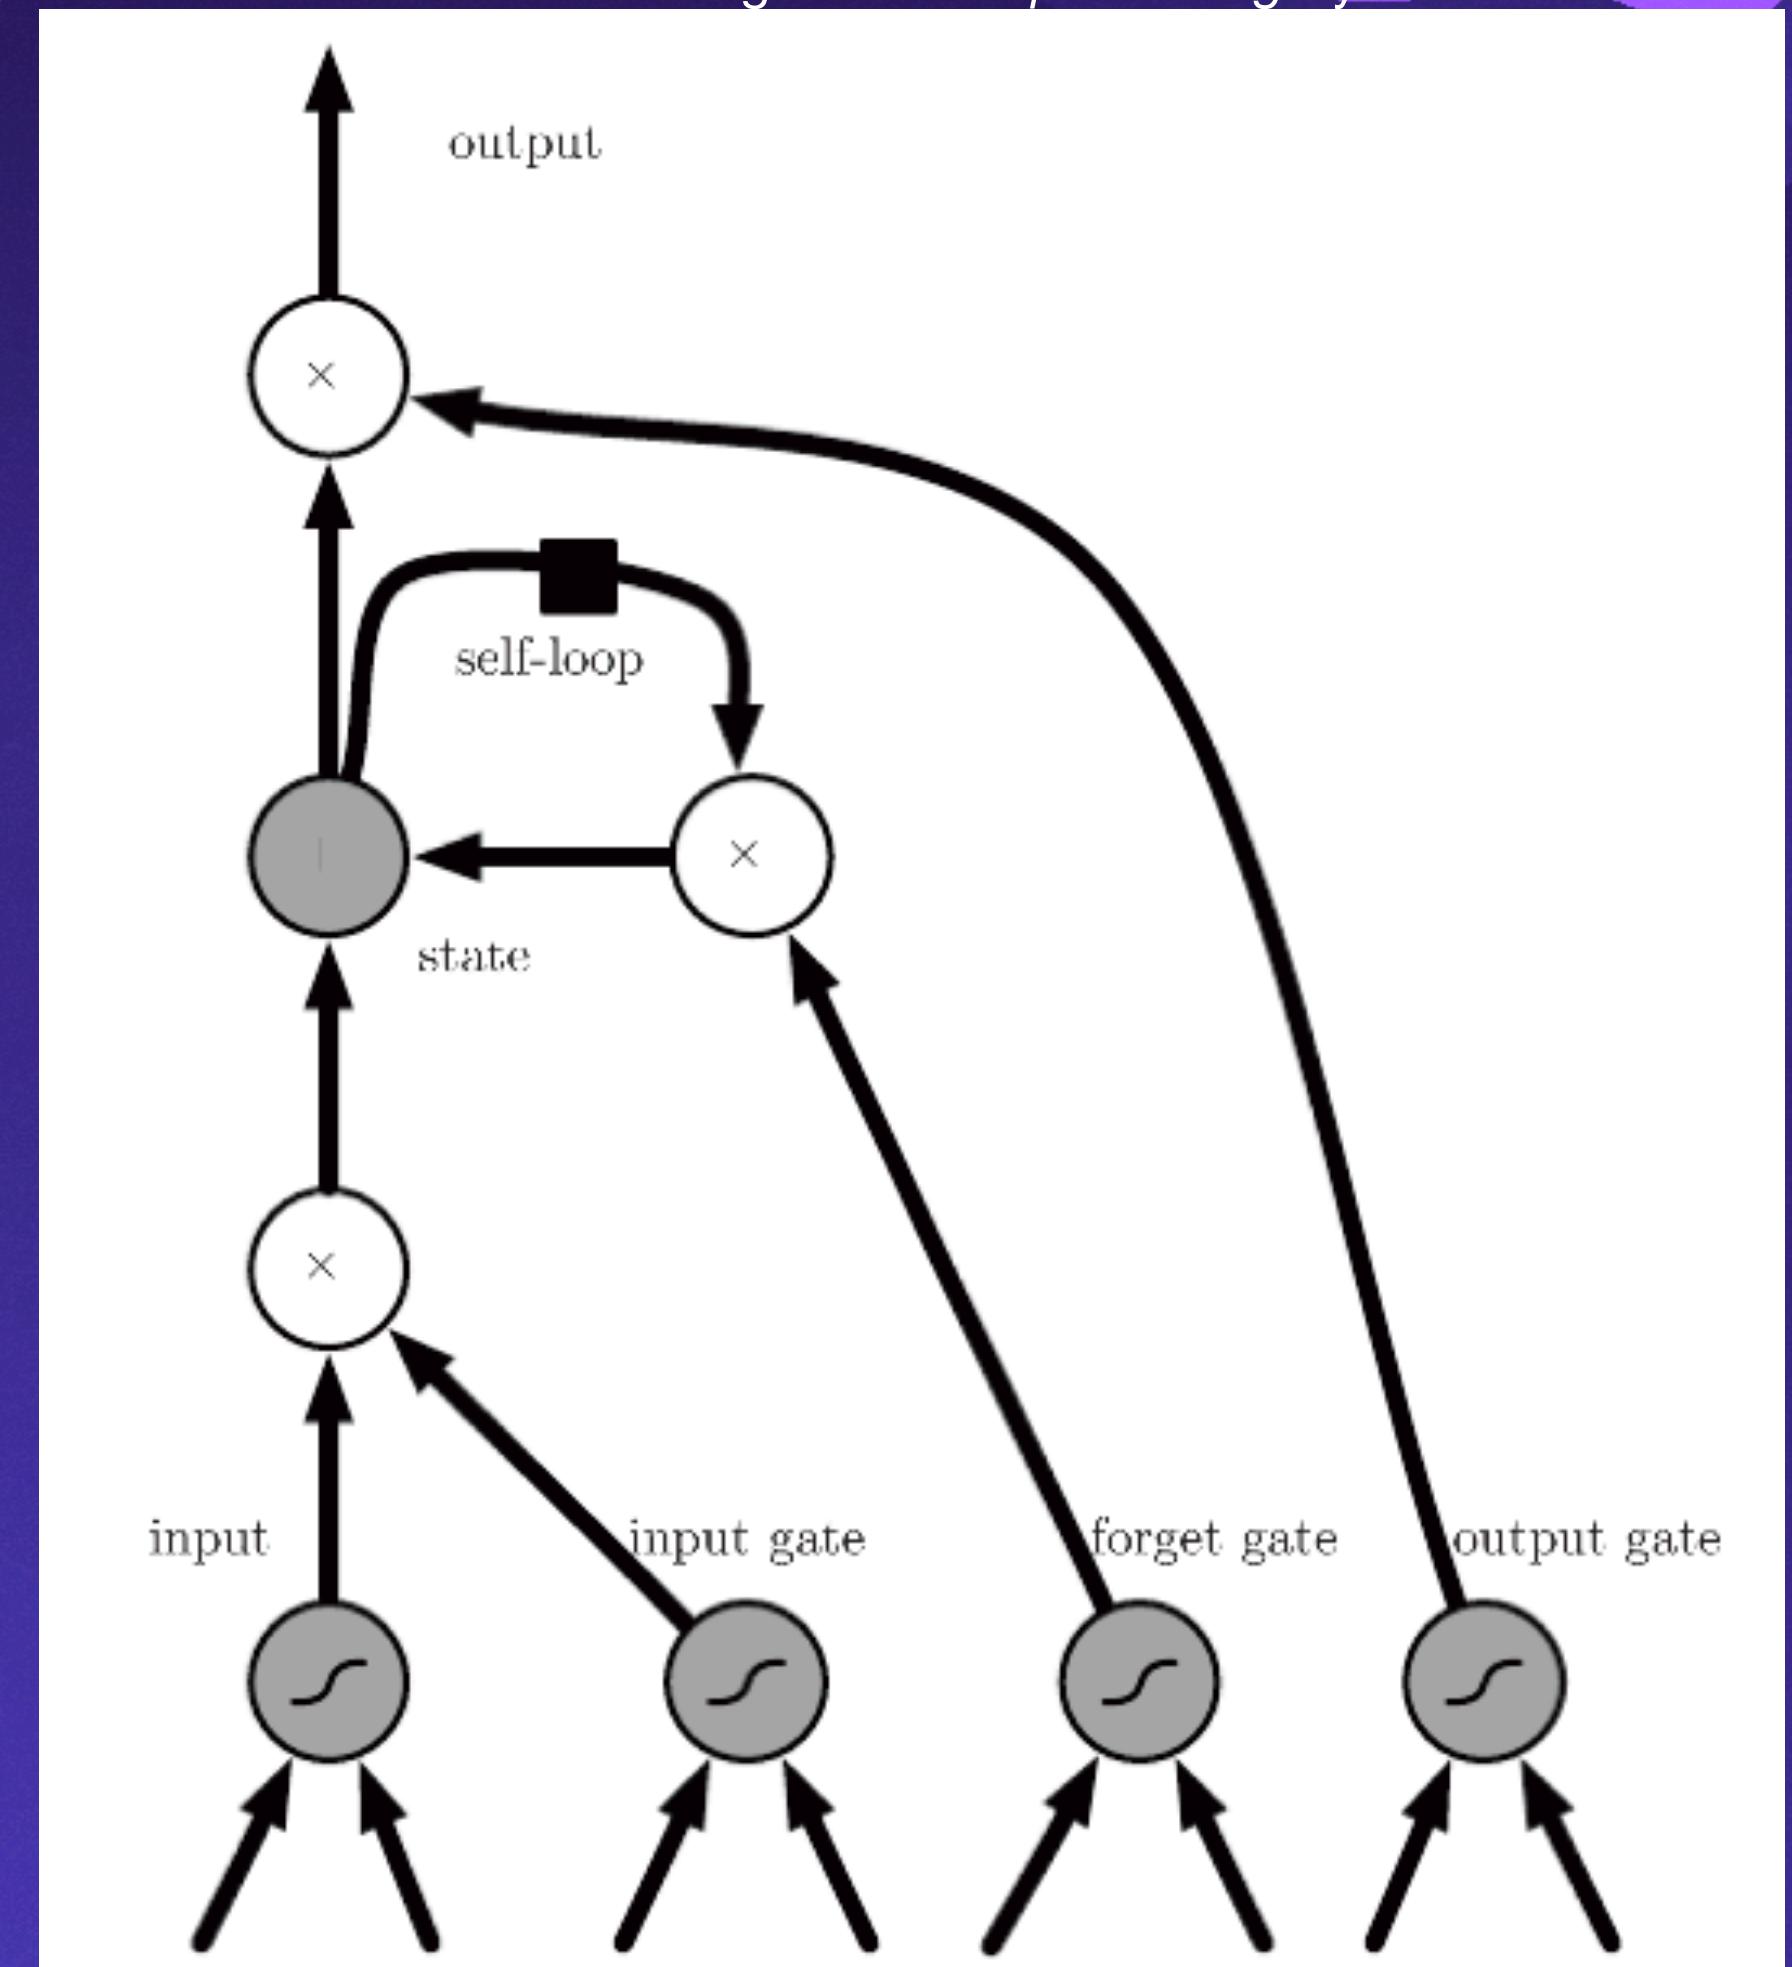
RNN with no outputs; information is processed sequentially, taking into account both $\mathbf{x}^{(t)}$ and $h^{(t-1)}$ but applying the *same function* $f(\cdot; \theta)$ at each timestep

Long Short-Term Memory (LSTM)

An Upgraded RNN Module

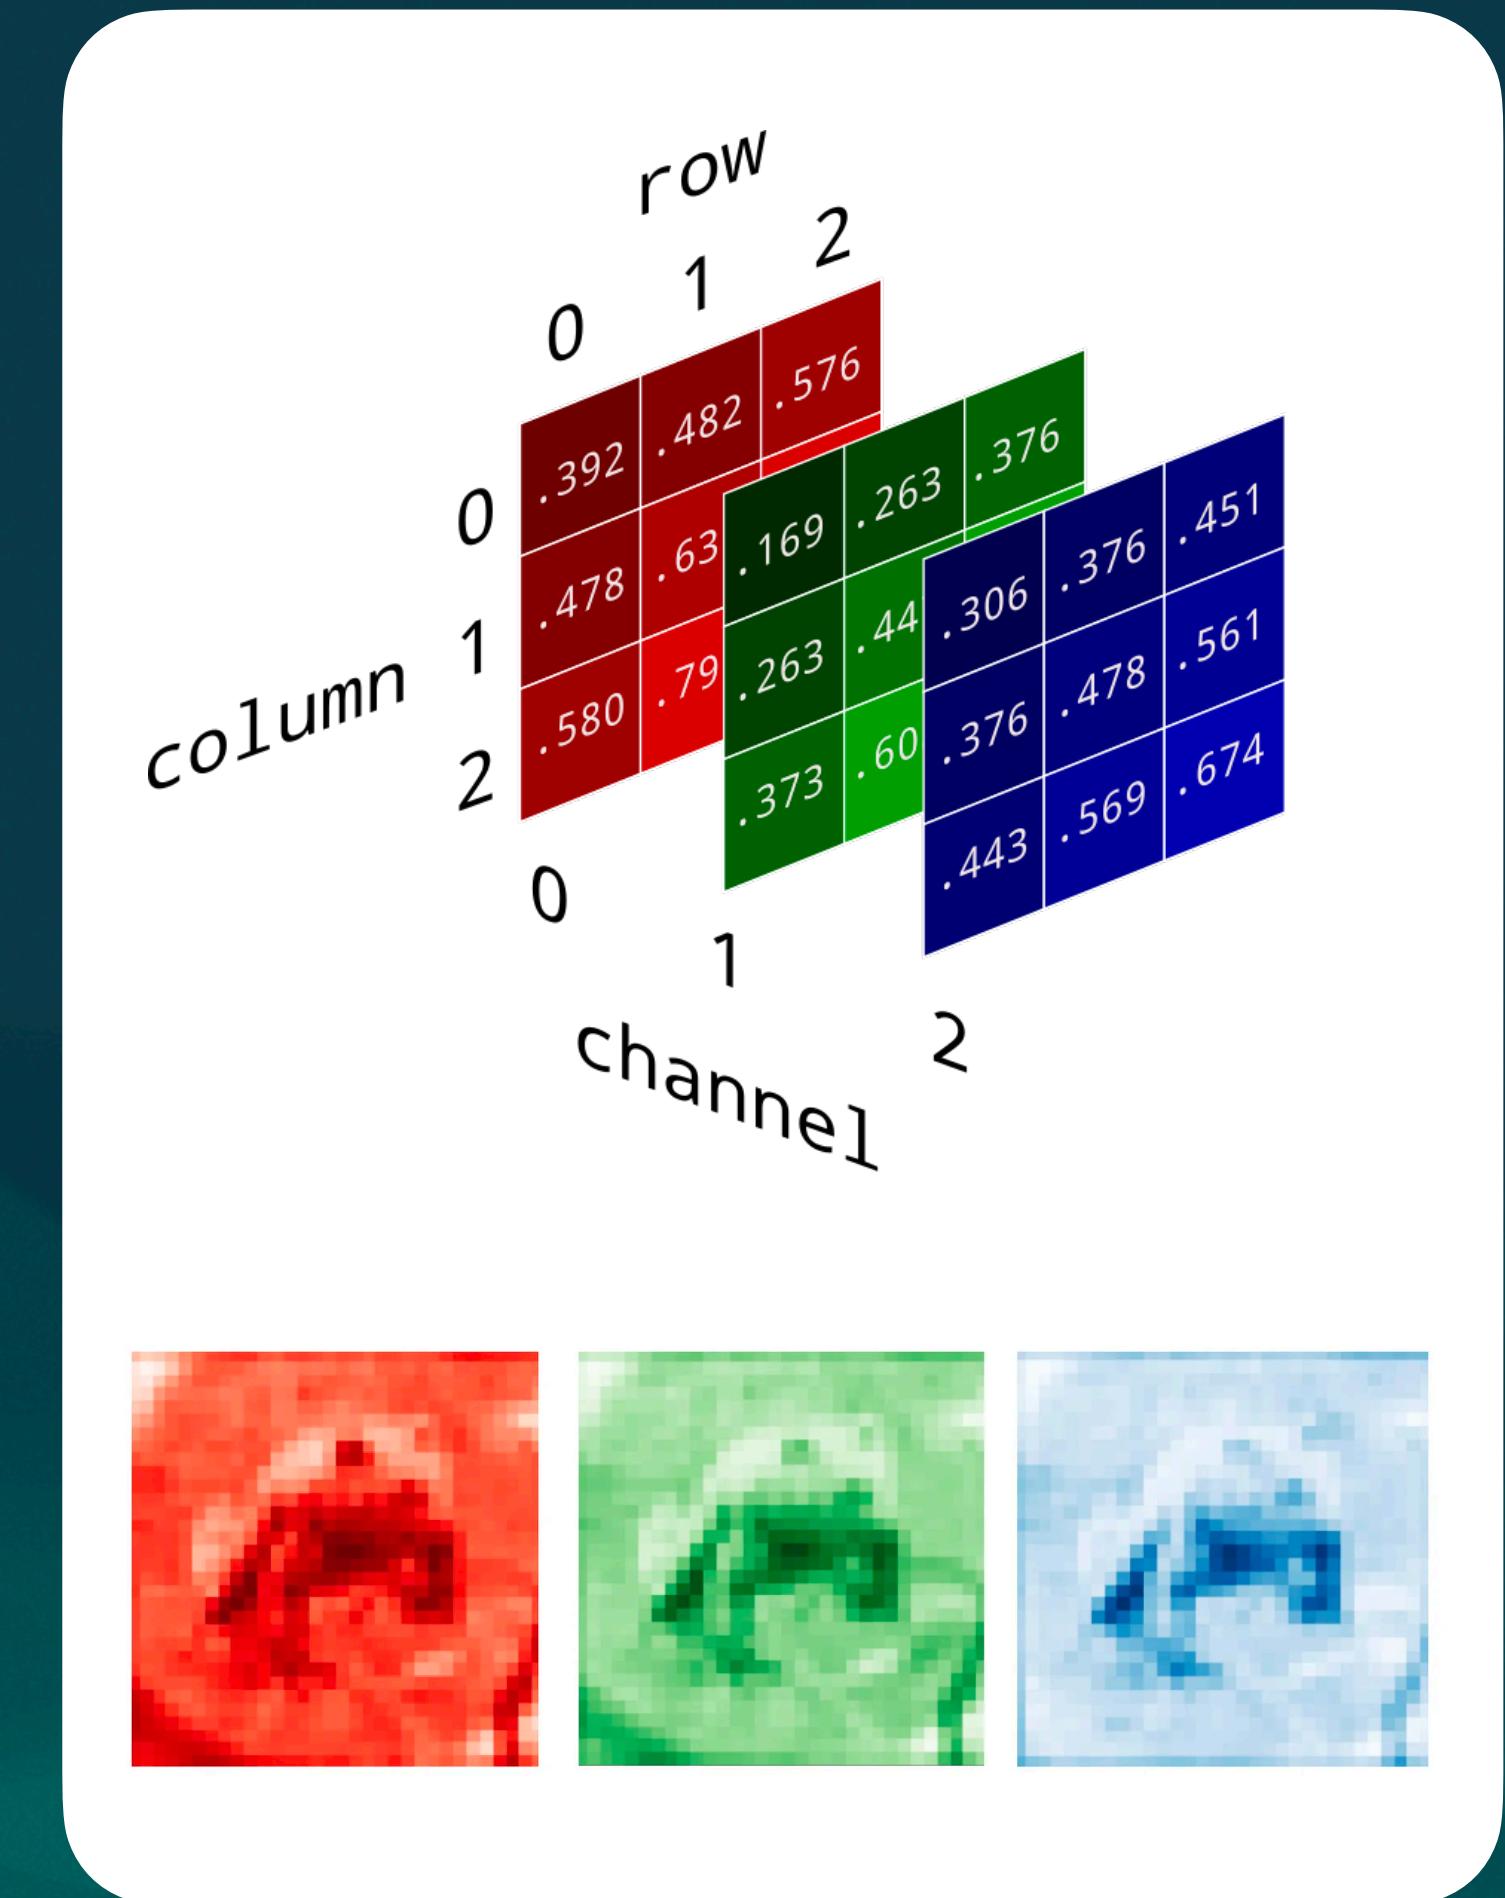
- RNNs are finicky to train; they often suffer from exploding/vanishing gradients
- This has motivated the development of more advanced RNNs like LSTMs
- The LSTM is a recurrent “cell” that is applied to all timesteps equally

Image from *Deep Learning* by Goodfellow et al.



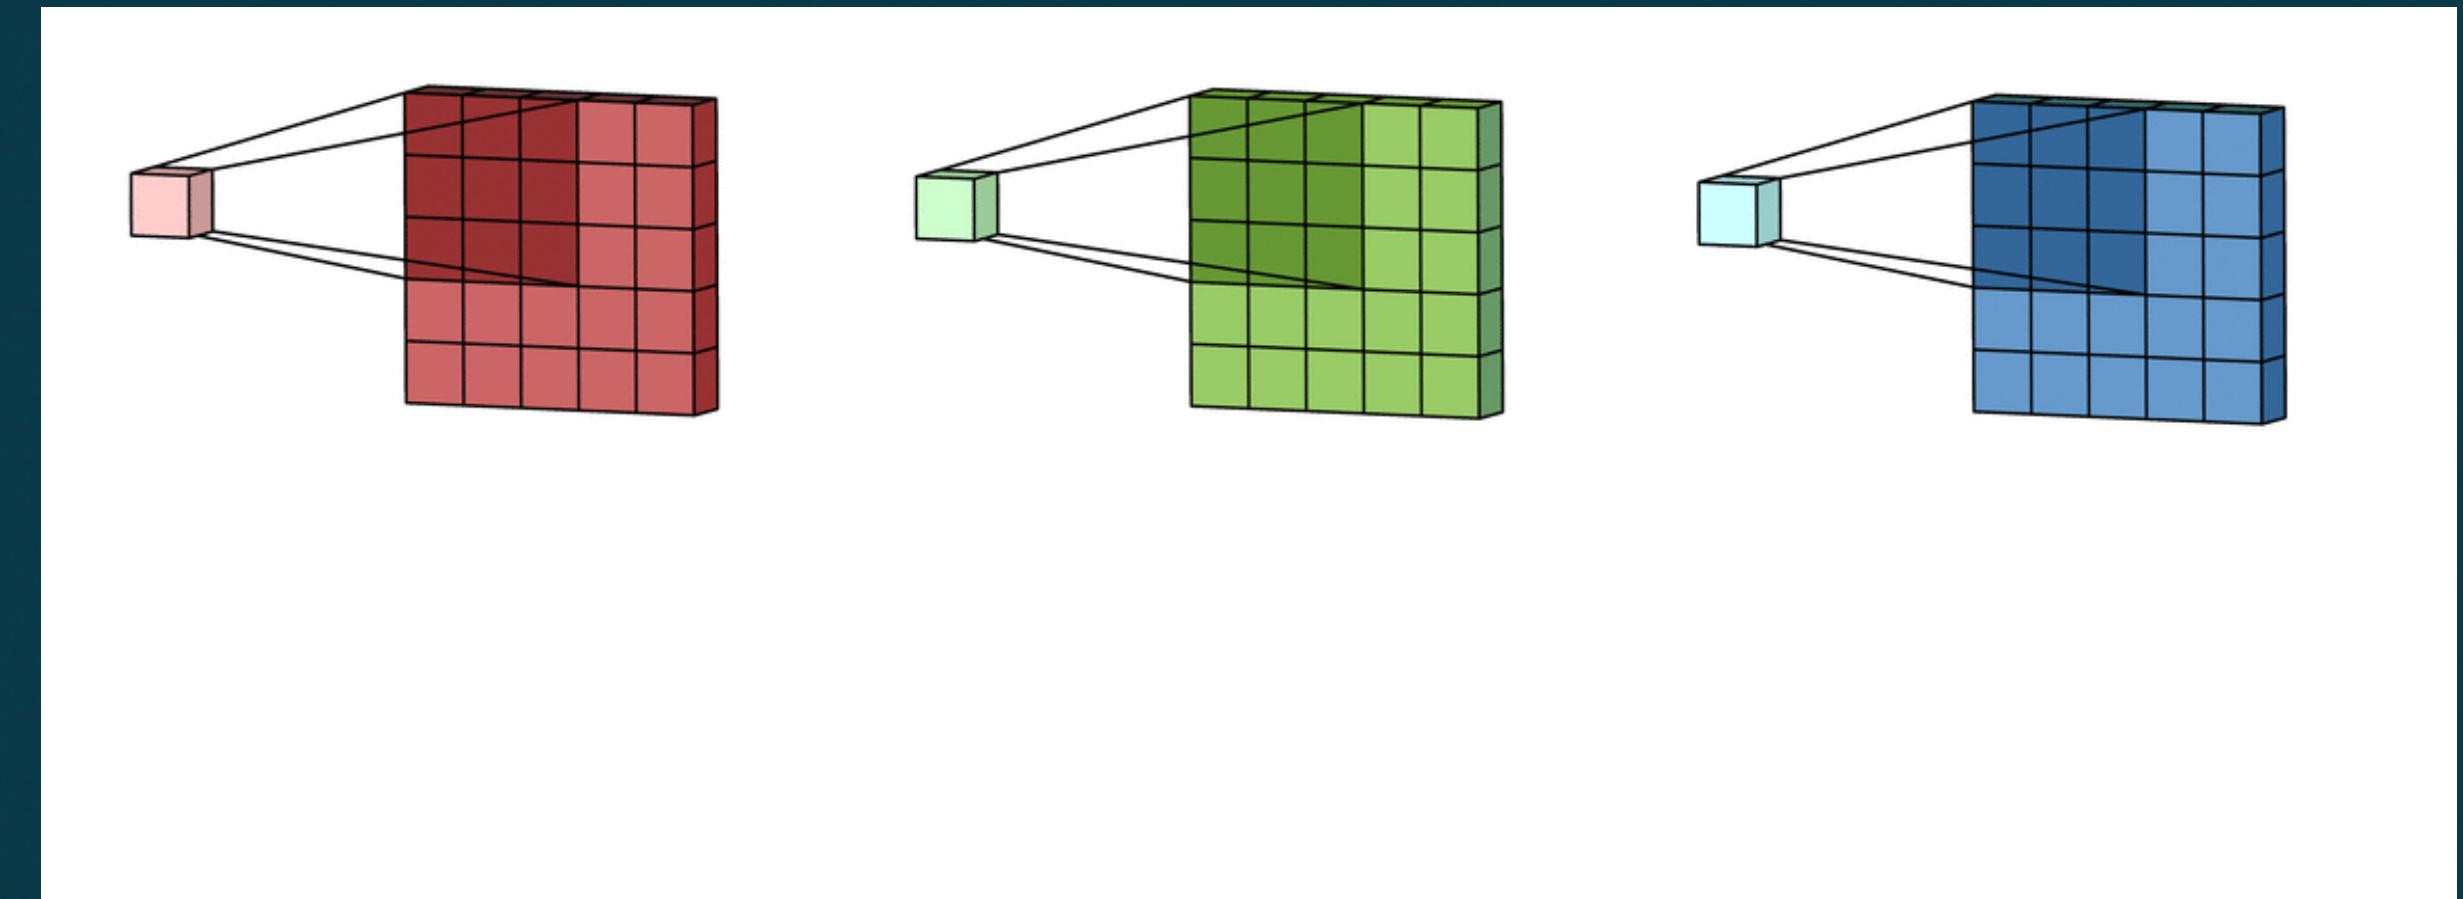
Convolutional Neural Networks (CNNs)

- Deep learning applied to images (data on a grid)
- For square images, inputs are $I \in \mathbb{R}^{n_{\text{pixels}} \times n_{\text{pixels}} \times n_{\text{features}}}$

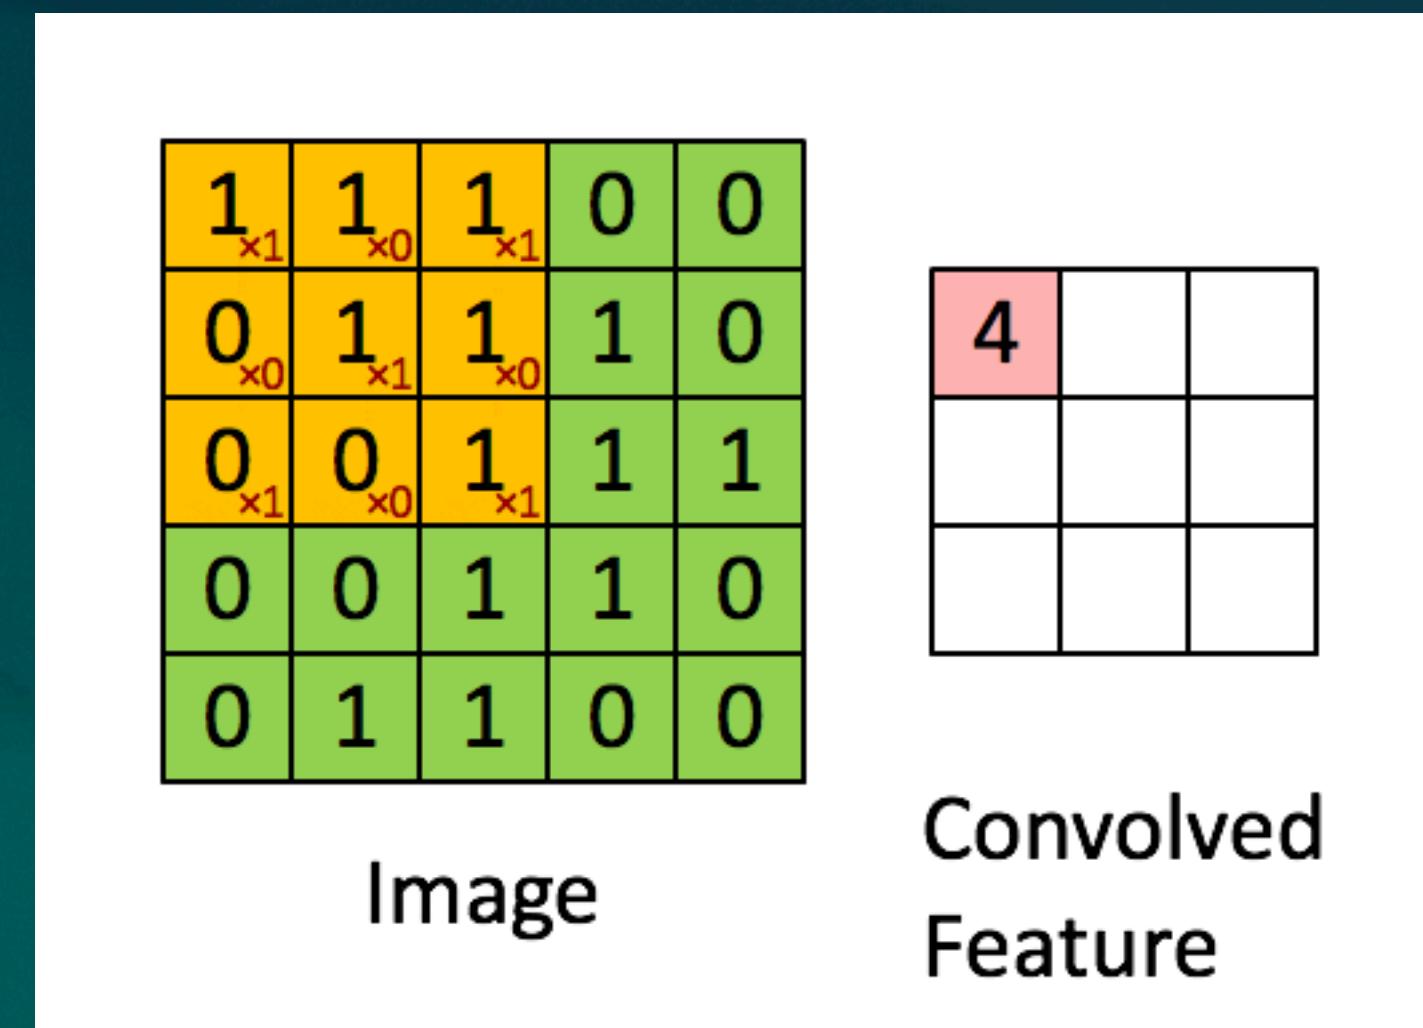


Convolutional Neural Networks (CNNs)

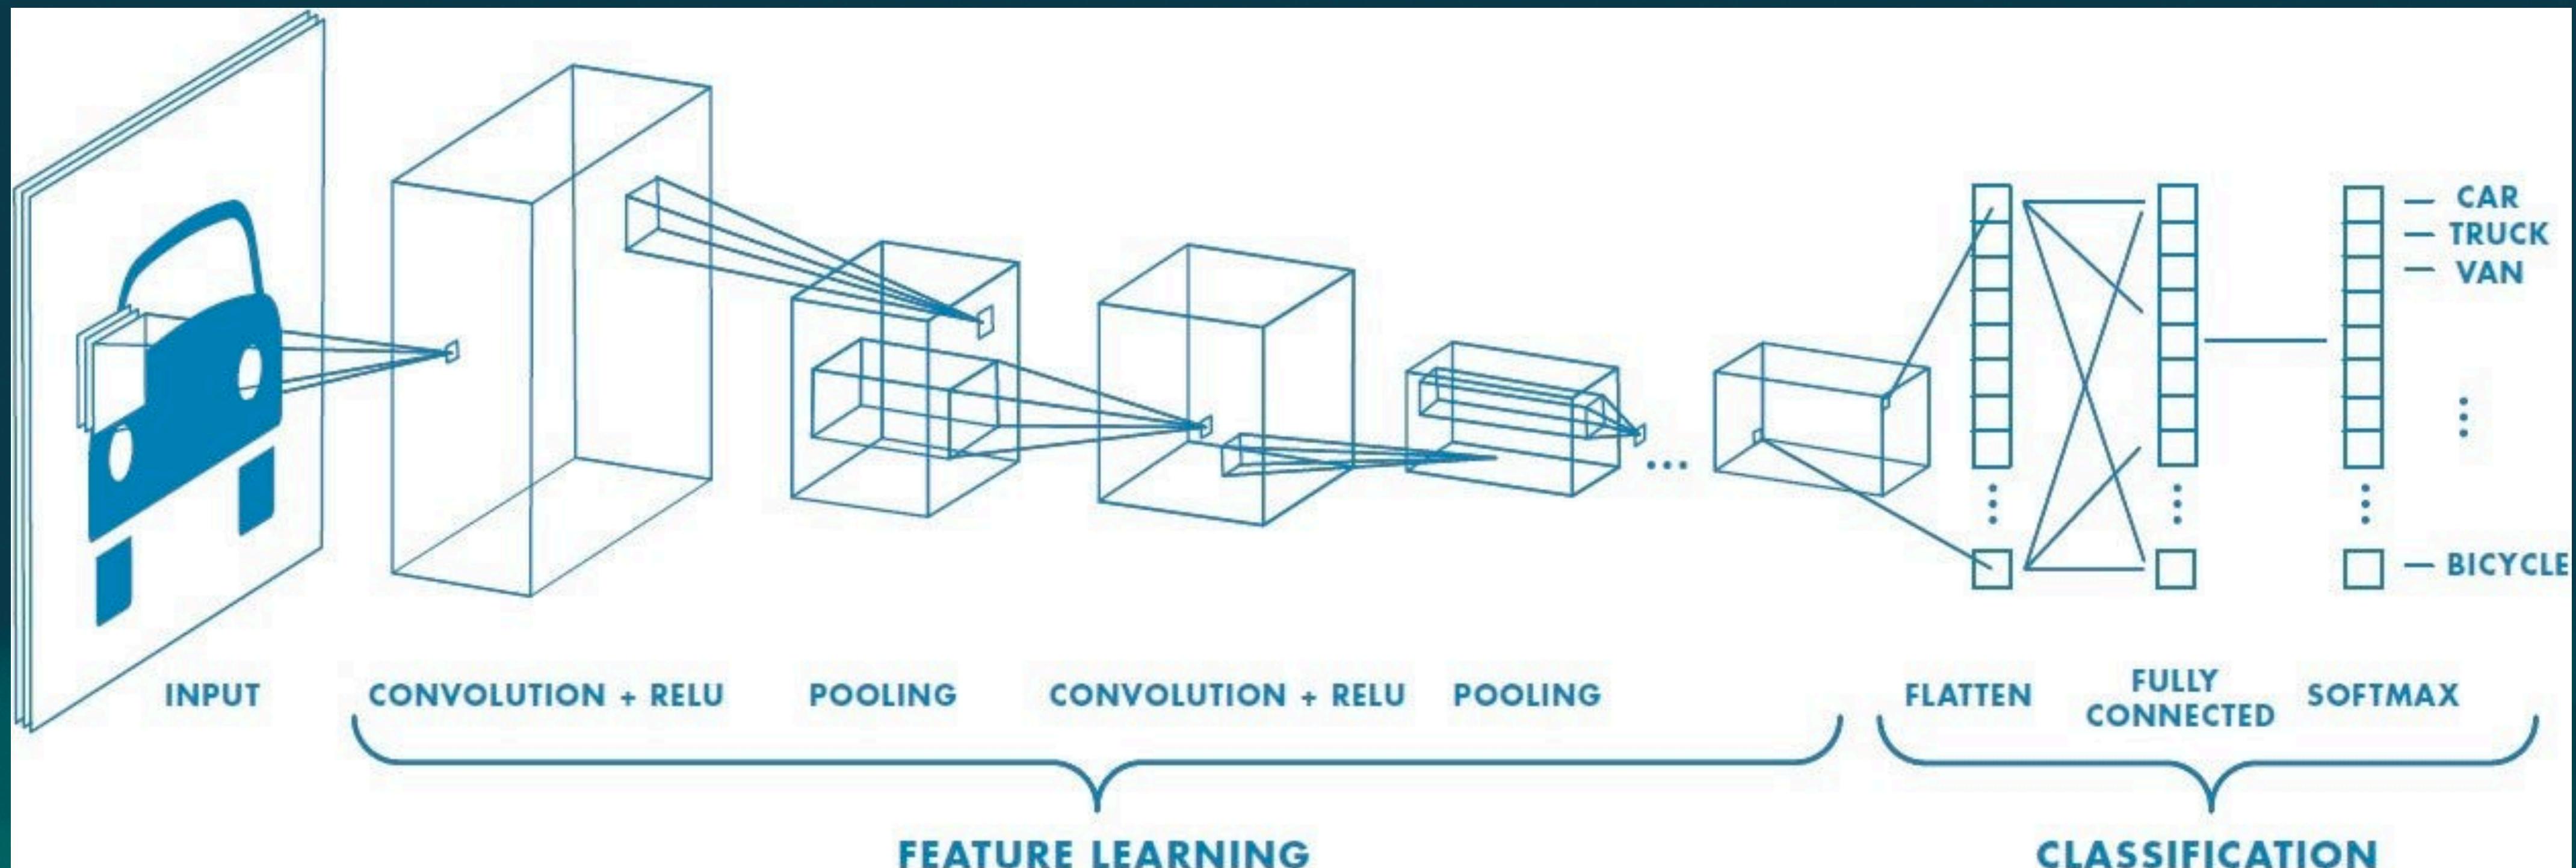
- Very similar approach to DNNs (non-sequential inputs, feed-forward approach), except now we use *convolutional* layers



- Convolution: filter is *convolved* (weighted sum with learnable weights) with the input image
- Again, parameter sharing!



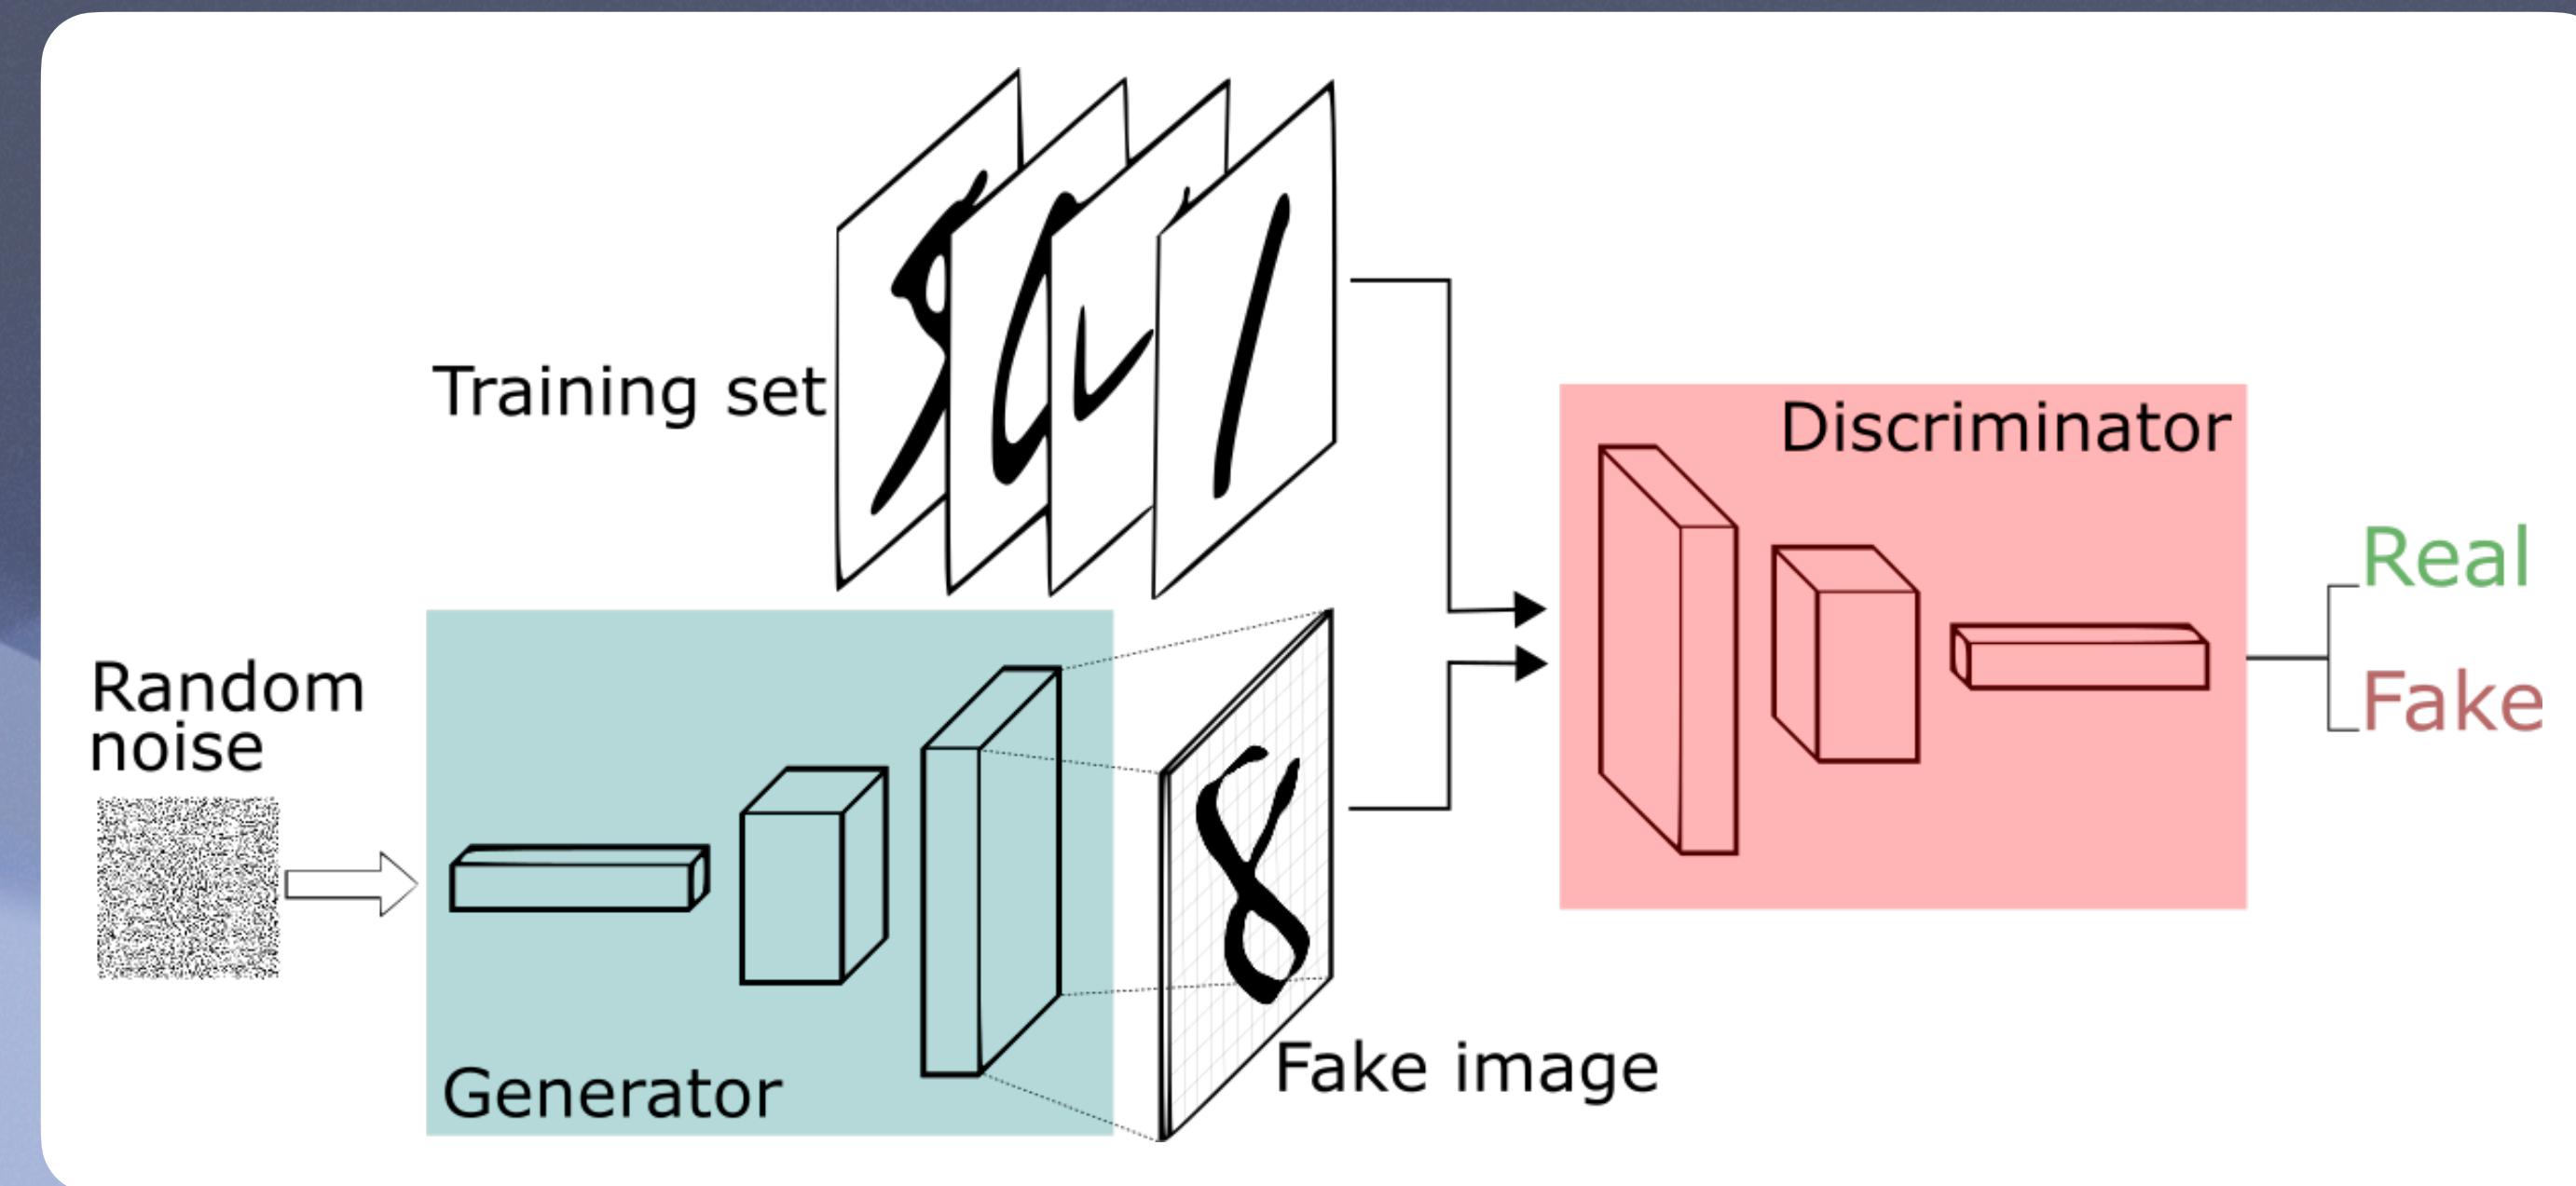
Convolutional Neural Networks (CNNs)



Generative Adversarial Networks (GANs)

Survey of Deep Learning Architectures

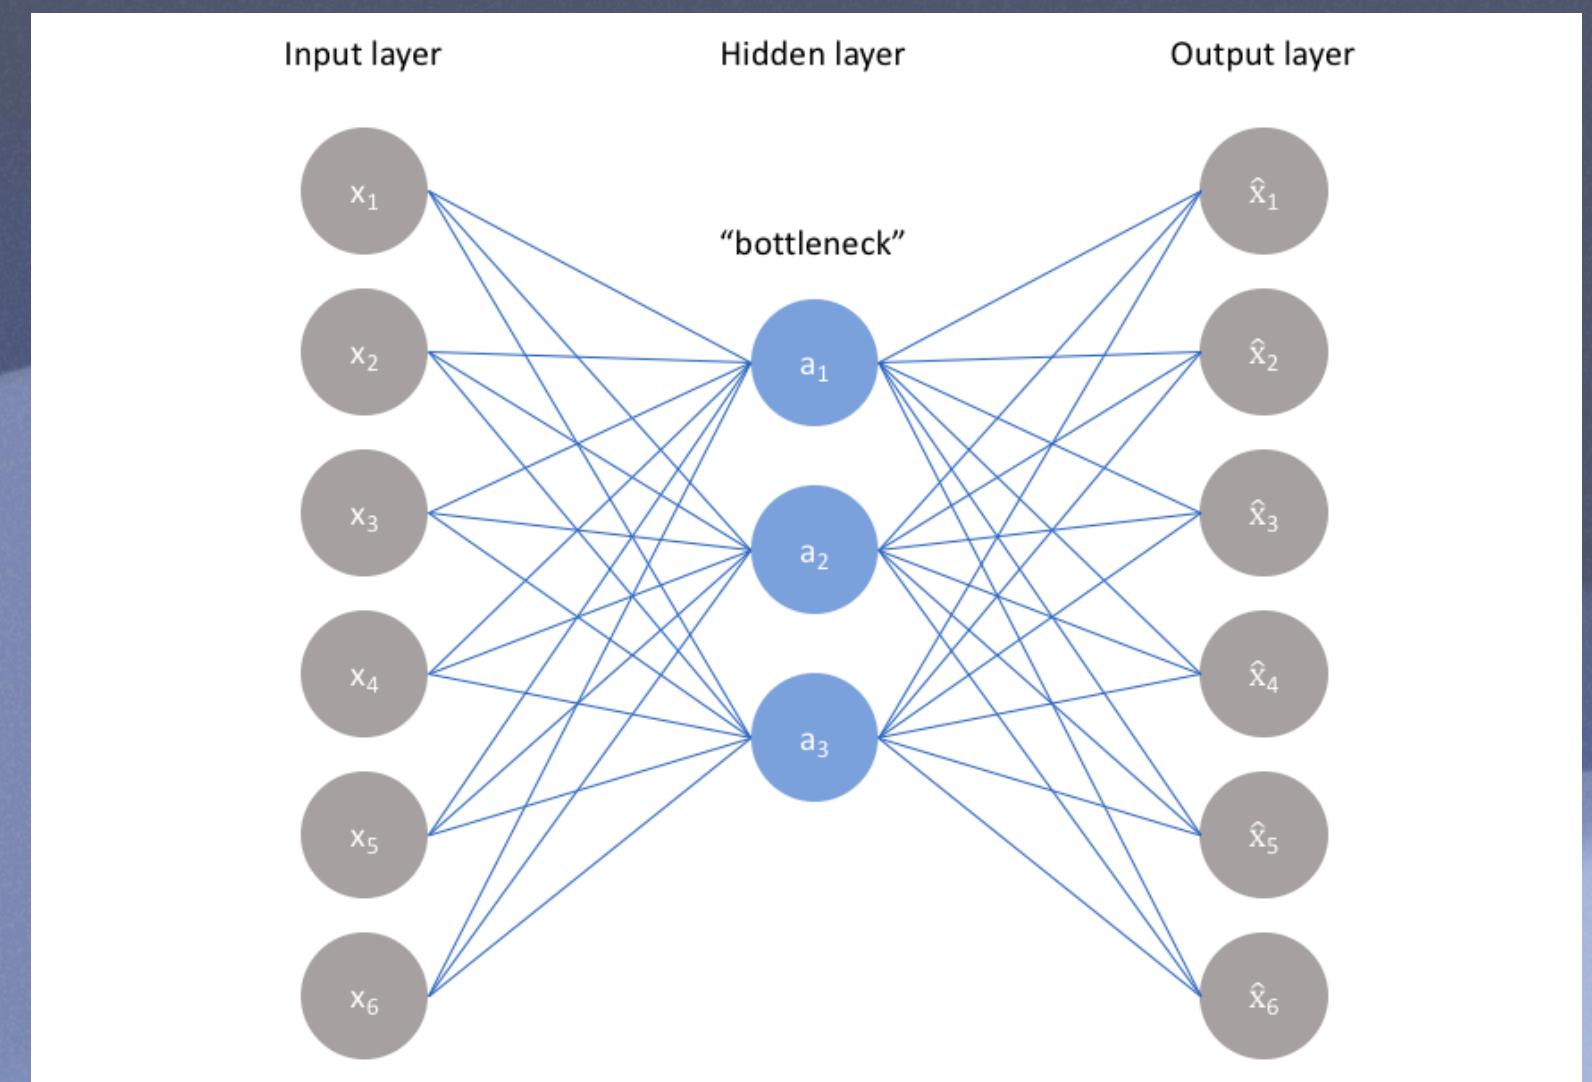
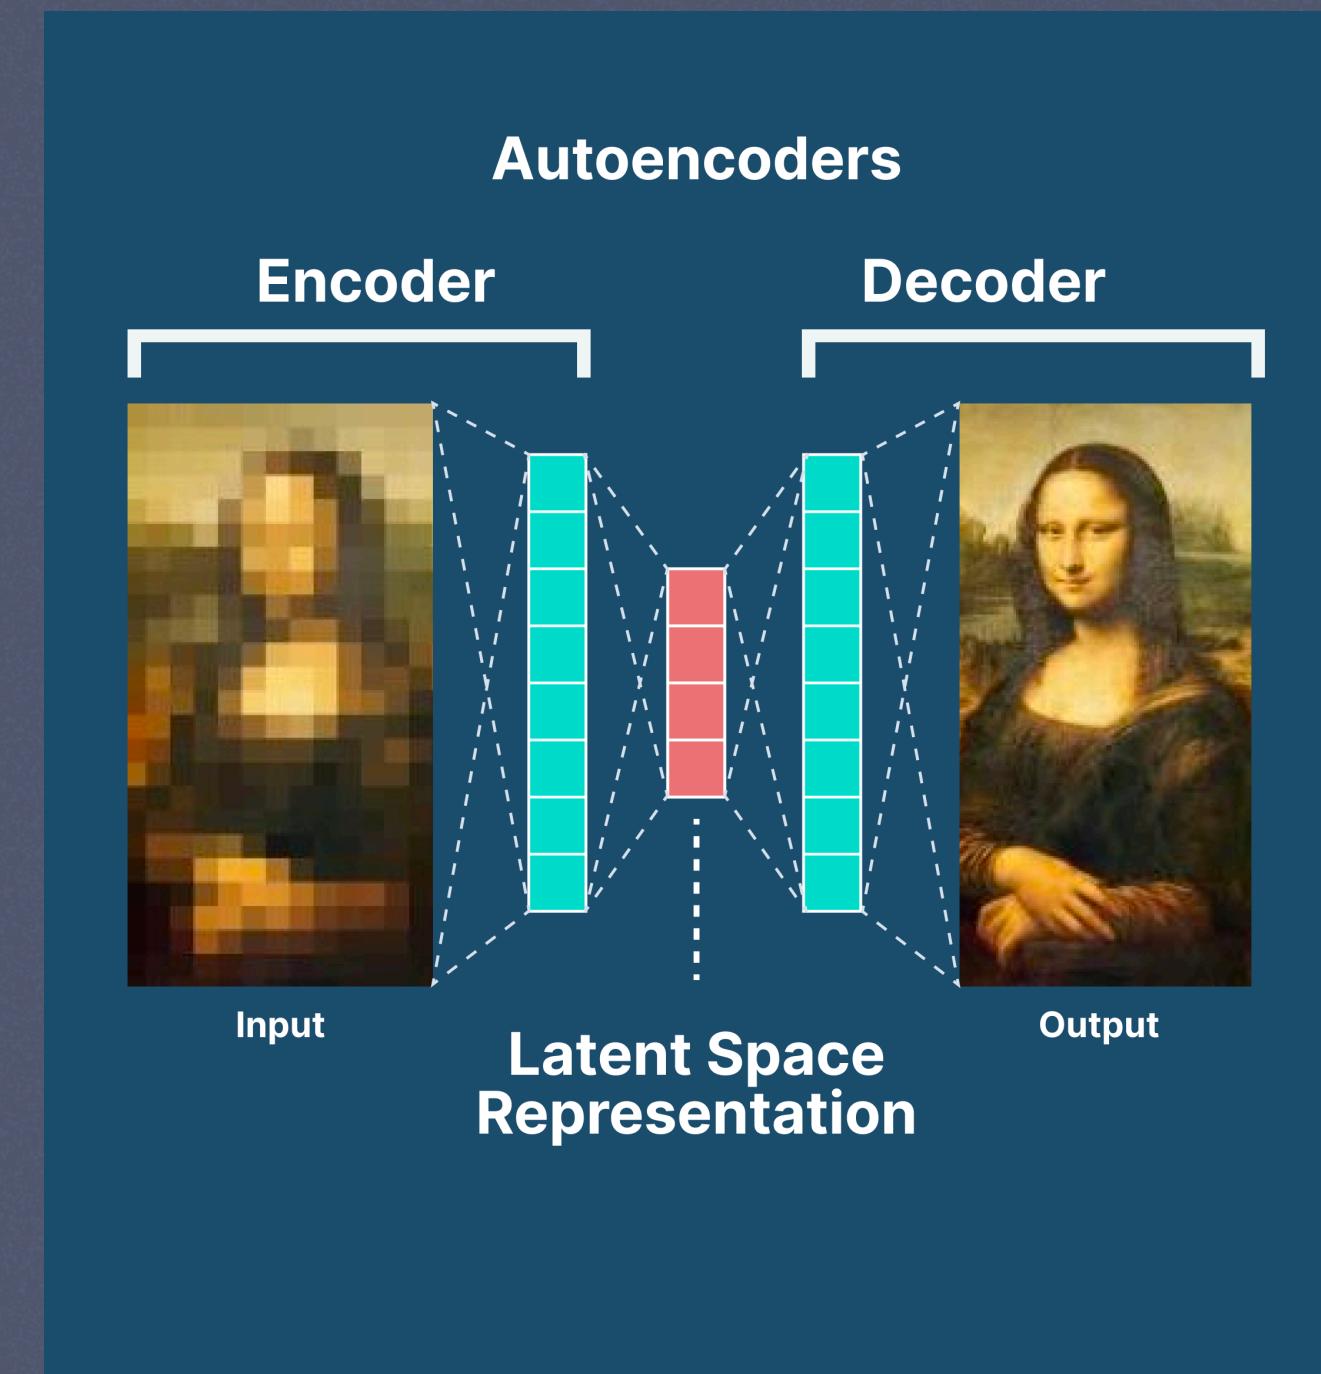
- “Generative” AI: use ML to create new images, sounds, etc.
- GANs: two agents (*the generator* and *the discriminator*) are given competing tasks:



Autoencoders

Learning Efficient Codings

- Autoencoders are used to produce compressed data representations
 - **Encoder:** produces a lower-dimensional (compressed) “latent” representation of the input data
 - **Decoder:** given the compressed representation, reconstruct the original data
- Decoded representations typically less noisy,
- Uses: efficient encoding, image denoising, generative modeling, anomaly detection



<https://www.v7labs.com/blog/autoencoders-guide#:~:text=An%20autoencoder%20is%20an%20unsupervised,even%20generation%20of%20image%20data.>

Variational Autoencoder (VAEs)

Generative Modeling via Autoencoders

- Generate realistic images from random noise
 - **Encoder:** predict means and standard deviations of a *probability distribution* over the latent features
 - **Decoder:** given a random sample from the latent distributions, produce the corresponding output

